

MASTERARBEIT | MASTER'S THESIS

Titel | Title

Termination results for hybrid approach of Joint Spectral Radius
computation

verfasst von | submitted by
Aaron Pumm, BA

angestrebter akademischer Grad | in partial fulfilment of the requirements for the degree of
Master Zusatz (Abkürzung)

Wien, | Vienna, 2025

Studienkennzahl lt. Studienblatt | degree
programme code as it appears on the student
record sheet:

UA 066 821

Studienrichtung lt. Studienblatt | degree
programme as it appears on the student record
sheet:

Masterstudium Mathematik

Betreut von | Supervisor:

Assoz. Prof. Dr. Martin Ehler

Contents

1	Introduction	5
1.1	Motivation of the JSR	5
1.2	Theoretical background	6
1.3	Preprocessing	8
2	Invariant-polytope algorithm	9
2.1	Extremal norms	9
2.2	Invariant polytope norms	9
2.3	Structure of the invariant-polytope algorithm	9
2.4	Stopping criterions	10
2.5	Termination conditions	10
2.6	Rebalancing and added starting vertices	10
2.7	Eigenvector cases	10
3	Finite-tree algorithm	11
3.1	Notation and definitions	11
3.2	Structure of the finite-tree algorithm	12
3.3	Efficiency results	13
4	Hybrid approach	15
4.1	Goal and options	15
4.2	Structure of the hybrid algorithm	15
4.3	Termination results	17
4.4	Efficiency	19
5	Numerics	20
5.1	Runtime of algorithm 3 on generated matrices	20
5.1.1	Variing matrix dimensions	20
5.1.2	Variing matrix condition	20
5.1.3	Variing matrix s.m.p number	20
5.2	Comparison algorithm 3 and 1	20
5.2.1	Big matrix dimensions	20
5.2.2	Multiple s.m.p's	20
5.2.3	old closed problems couldnt be solved by 1	20
5.3	Comparison algorithm 3 and 2	20
5.3.1	Multiple s.m.p's and big matrix dimensions	20
5.3.2	old closed problems couldnt be solved by 2	20

5.4	Solving open problems	20
6	Conclusion	21
6.1	New approach	21
6.2	Solved Problems	21
6.3	Efficiency	21
6.4	Unsolved problems	21
6.5	Further research	21
6.5.1	Complex case	21
6.5.2	Nonnegative case	21
6.5.3	Optimizing tree generation and vertex choice	21

Zusammenfassung

Diese Arbeit präsentiert einen hybriden Ansatz zur Berechnung des gemeinsamen Spektralradius (JSR) einer endlichen Menge von Matrizen. Die vorgeschlagene Methode integriert die invariant-polytope und finite-tree Algorithmen, um deren jeweilige Stärken zu nutzen. Die breite Anwendbarkeit des baumbasierten Ansatzes wird mit der Effizienz der Polytop-basierten Techniken kombiniert. Der entwickelte Algorithmus konstruiert eine strukturierte Baumdarstellung und überprüft maximale Produkte mithilfe angepasster Normen, wodurch eine Terminierung in Fällen ermöglicht wird, in denen klassische Methoden an ihre Grenzen stoßen (Bedingungen an Terminierung oder Laufzeit). Theoretische Garantien werden bereitgestellt, um Korrektheit und Effizienz zu belegen. Es werden auch JSR-Berechnungen für Matrizenmengen, bei denen bestehende Methoden scheitern oder ineffizient sind demonstriert.

Abstract

This thesis presents a hybrid approach for computing the joint spectral radius (JSR) of a finite set of matrices. The proposed method integrates the invariant polytope algorithm and the finite tree algorithm to leverage their respective strengths—combining the broad applicability of the tree-based approach with the efficiency of polytope-based techniques. The developed algorithm constructs a structured tree representation and verifies spectrum-maximizing products using adapted norms, enabling termination in cases where classical methods struggle. Theoretical guarantees are provided to demonstrate correctness and efficiency. Applications include improved JSR computations for matrix sets where existing methods fail or prove inefficient.

Keywords: Joint spectral radius, hybrid algorithm, invariant polytope, finite tree method, matrix norms.

1 Introduction

The *Joint Spectral Radius* (JSR) was first introduced by G.-C. Rota and G. Strang in 1960 (Rota and Gilbert Strang, 1960). They described the JSR as the maximal exponential growth rate of a product of matrices taken from a finite set. Since its inception, the JSR has become a cornerstone in various mathematical and engineering disciplines due to its ability to encapsulate the asymptotic behavior of matrix sequences.

The concept gained significant traction in the 1990s when researchers began exploring its theoretical properties and practical implications. Notable advancements include its application in control theory, where it is used to analyze the stability of switched linear systems (Blondel and Tsitsiklis, 2000), and in wavelet theory, where it assists in the construction of refinable functions (citepdaubechies1992ten). The computational challenges associated with determining the JSR have inspired the development of several algorithms, such as the invariant-polytope method (Guglielmi and Protasov, 2013) and the finite-tree method (Jungers, 2009).

Despite the progress, the JSR computation remains a challenging problem, particularly due to the exponential complexity of exploring all possible matrix products. This thesis seeks to contribute to this ongoing effort by leveraging the invariant-polytope algorithm and the finite-tree algorithm to create a hybrid methodology that mitigates their respective limitations.

To fully grasp the subsequent mathematical framework, the reader should be familiar with linear algebra, specifically matrix norms, eigenvalues, and spectral radius. A basic understanding of combinatorial optimization and algorithm design will also be beneficial.

1.1 Motivation of the JSR

For some motivation i would like to point our interest to the analysis of wavelets:

Structure of the Thesis

The remainder of this thesis is structured as follows: Chapter 1 provides a sufficient background on the JSR and its basic properties. Chapters 2 and 3 present the ideas and concepts of the algorithms that will be exploited to form the proposed hybrid approach, outlining their theoretical foundation and algorithmic implementation. Chapter 4 discusses possible combinations of former approaches, proposes the so-called Tree-flavored-invariant-polytope algorithm, and brings proofs of termination which is the main result of this thesis. *Chapter 5* presents numerical results on ... problems to analyze the efficiency and applicability of the hybrid algorithm. Chapter 6 concludes with insights and future directions.

1.2 Theoretical background

The *joint spectral radius* (JSR) of a set of matrices is a generalization of the spectral radius for a single matrix. For a finite set of matrices $\mathcal{A} = \{A_1, A_2, \dots, A_m\}$, the JSR is formally defined as:

$$JSR(\mathcal{A}) = \lim_{k \rightarrow \infty} \max_{A_i \in \mathcal{A}} \|A_{i_k} \dots A_{i_1}\|^{1/k}, \quad (1.1)$$

where $\|\cdot\|$ denotes any submultiplicative matrix norm.

Proof. well-definedness

Let $\|\cdot\|_1$ and $\|\cdot\|_2$ be two submultiplicative norms on $\mathbb{R}^{n \times n}$. By equivalence of norms in finite-dimensional vector spaces, there exist constants $c, C > 0$ such that:

$$c\|P\|_1 \leq \|P\|_2 \leq C\|P\|_1 \quad \forall P \in \mathbb{R}^{n \times n}$$

Now if we consider this and take asymptotic equality into account we get:

$$\begin{aligned} & \lim_{k \rightarrow \infty} \max_{A_i \in \mathcal{A}} \|A_{i_k} \dots A_{i_1}\|_1^{1/k} \\ & \leq \lim_{k \rightarrow \infty} \left(\frac{1}{c}\right)^{\frac{1}{k}} \max_{A_i \in \mathcal{A}} \|A_{i_k} \dots A_{i_1}\|_2^{1/k} \\ & = \lim_{k \rightarrow \infty} \max_{A_i \in \mathcal{A}} \|A_{i_k} \dots A_{i_1}\|_2^{1/k} \\ & \leq \lim_{k \rightarrow \infty} C^{\frac{1}{k}} \max_{A_i \in \mathcal{A}} \|A_{i_k} \dots A_{i_1}\|_1^{1/k} \\ & = \lim_{k \rightarrow \infty} \max_{A_i \in \mathcal{A}} \|A_{i_k} \dots A_{i_1}\|_1^{1/k} \end{aligned}$$

□

A simple application would be the categorization of the stability of linear switched dynamical systems in discrete time. Let's define $v_{n+1} := A_{i_n} v_n$ where $i_n \in \{1, \dots, m\}$. Now the system is stable (i.e. $v_n \rightarrow 0$) exactly if $JSR(\mathcal{A}) < 1$.

This hints at the more general problem of finding whether $JSR(\mathcal{A}) < 1$ which is proven to be mathematically undecidable [Jungers] as seen in a future section. This and some other complexity results are the reason of vast research efforts in JSR analysis and numerics.

To understand the state-of-the-art algorithms considered as well as the main results that will follow, it is necessary to introduce some basic properties of the JSR.

Homogeneity

For any scalar α and set of matrices \mathcal{A} , the scaling property

$$JSR(\alpha\mathcal{A}) = |\alpha| JSR(\mathcal{A}) \quad (1.2)$$

holds, which follows directly from norm homogeneity.

Irreducibility

A set of matrices is called (*commonly*) *reducible* if there exists a nontrivial subspace of \mathbb{R}^n that is invariant under all matrices in the set. This means there exists a change of basis that block-triangularizes all matrices in \mathcal{A} at the same time. If \mathcal{A} is not reducible it is called irreducible.

Three-member inequality

The *three-member inequality* provides essential bounds for the JSR. For any submultiplicative matrix norm $\|\cdot\|$, the inequality

$$\max_{P \in \mathcal{A}^k} \rho(P)^{\frac{1}{k}} \leq JSR(\mathcal{A}) \leq \max_{P \in \mathcal{A}^k} \|P\|^{\frac{1}{k}}, \quad (1.3)$$

holds for every $k \in \mathbb{N}$ (Jungers, 2009). This result forms a starting point for many computational approaches as the bounds are sharp in the sense that both sides converge to the JSR as $k \rightarrow \infty$ (left side in \limsup).

Minimum over norms

The JSR can be equivalently defined as the minimum over all submultiplicative norms:

$$JSR(\mathcal{A}) = \min_{\|\cdot\|} \max_{A \in \mathcal{A}} \|A\|. \quad (1.4)$$

Proof. equivalence

□

Complexity and the Finiteness Property

The computation of the JSR is known to be computationally challenging, with determining its exact value classified as NP-hard. However, certain matrix sets exhibit the *finiteness property*, which states that there exists a finite sequence of matrices A_{i_1}, \dots, A_{i_k} such that:

$$JSR(\mathcal{A}) = \|A_{i_k} \cdots A_{i_1}\|^{1/k}. \quad (1.5)$$

While this property does not hold universally, it is essential for algorithmic approaches.

Candidates and Generators

Approximating the JSR requires identifying candidate products or *generators* of the matrix set that contribute most significantly to the asymptotic growth rate. These generators are often derived through optimization techniques and their identification is a key step in computational algorithms.

Similarity and reducibility

In the following section it is assumed that the set of interest is irreducible for some later results. If the set happens to be reducible there is an easy work-around but for that we need an intermediate result: For any invertible matrix T and reducible set \mathcal{A}

$$JSR(\mathcal{A}) = JSR(T^{-1}\mathcal{A}T) \quad (1.6)$$

holds. Now per definition there exists a change of basis such that all $A \in \mathcal{A}$ are block-triangularized:

$$\exists T : T^{-1}A_iT = \begin{bmatrix} B_i & C_i \\ 0 & D_i \end{bmatrix}$$

Now

$$JSR(\mathcal{A}) = \max\{JSR(\{B_i\}), JSR(\{D_i\})\} \quad (1.7)$$

The proof can be seen in [Jungers]. This can be applied iteratively until the sets of blocks are all irreducible. The problem was split into similar problems of smaller dimension. For the following considerations we can now assume \mathcal{A} to be irreducible.

1.3 Preprocessing

This thesis aims to address the challenge of computing the JSR by combining two existing algorithms that have demonstrated practical effectiveness in calculating the JSR for nontrivial sets of matrices. Both algorithms are based on the following simple concept:

We want to find the JSR of the finite set of matrices $\mathcal{A} = \{A_1, \dots, A_n\}$

1. **Assumptions:** \mathcal{A} is irreducible and possesses the finiteness property.
2. **Candidates:** Efficiently find products $P = A_{i_k} \cdots A_{i_1}$ of matrices from \mathcal{A} that maximize the averaged-spectral radius $\hat{\rho} := \rho(P)^{\frac{1}{k}}$ for a given maximal length k_{\max} .
3. **Rescaling:** Transform $\mathcal{A} \rightarrow \tilde{\mathcal{A}}$ with $\tilde{A}_i := \frac{1}{\hat{\rho}} A_i$.
4. **Proofing:** Now establish the fact that $JSR(\tilde{\mathcal{A}}) = 1$ using the three-member-inequality. By homogeneity this is equivalent to $JSR(\mathcal{A}) = \hat{\rho}$.

The considered algorithms only differ in step 4, while the invariant-polytope algorithm tries to find a norm that bounds the products of length 1 already enough. The finite tree algorithm, on the other hand, bounds the products using some partitioning-space that separates every product into products that are 1-bounded and some rest-term that doesn't grow fast enough to overcome the k -th root of the JSR definition (polynomial growth). By integrating these algorithms into a hybrid approach, this work aims to advance the computational tools available for JSR analysis combining efficiency and a wide solution-space of the former. Where solution space is referring to the space of matrix-sets \mathcal{A} for which the algorithm finds the exact $JSR(\mathcal{A})$, not to be confused with the space of solutions like with differential equations.

2 Invariant-polytope algorithm

In this chapter we bring our interest to the underlying invariant-polytope algorithm. One result about JSR computation, that every irreducible family possesses an invariant norm is helpful. We observe that there always exists a norm that is in some sense extremal.

2.1 Extremal norms

A norm is called invariant if it is preserved under the action of the family of matrices. That is, for all $A_j \in \mathcal{A}$ and $x \in \mathbb{R}^d$ we have $\|A_j x\| = \|x\|$.

Theorem 2.1.1. *Every irreducible family \mathcal{A} possesses an invariant norm.*

Definition 2.1.2. *A norm $\|\cdot\|$ is called extremal for a family of matrices \mathcal{A} if $\|A_j x\| \leq \rho(\mathcal{A})\|x\|$ for all $x \in \mathbb{R}^d$ and $A_j \in \mathcal{A}$.*

Every invariant norm is extremal

2.2 Invariant polytope norms

For every polytope there exists a *Minkowski norm*

2.3 Structure of the invariant-polytope algorithm

After the preprocessing is over we start with the leading eigenvector of the calculated candidate. Now every product with matrices from \mathcal{A} and vertices from V is added if the current polytope-norm is bigger or equal to 1.

Algorithm 1 invariant-polytope algorithm

```

 $V := \{v_1, \dots, v_M\}$ 
 $V_{\text{new}} \leftarrow V$ 
while  $V_{\text{new}} \neq \emptyset$  do
     $V_{\text{rem}} \leftarrow V_{\text{new}}$ 
     $V_{\text{new}} \leftarrow \emptyset$ 
    for  $v \in V_{\text{rem}}$  do
        for  $A \in \mathcal{A}$  do
            if  $\|Av\|_{\text{cos}(V)} \geq 1$  then
                 $V \leftarrow V \cup Av$ 
                 $V_{\text{new}} \leftarrow V_{\text{new}} \cup Av$ 
return  $\text{co}_s(V)$ 

```

2.4 Stopping criterions

The runtime of the algorithm 1 is not finite. Suitable conditions for stopping or recalculating a candidate have to be put in place for the algorithm to have a finite runtime. Of course a bare minimum of an max iteration trigger is implemented. The paper (Guglielmi and Protasov, 2011) promises at least good bounds on the real JSR value while also proposing a stopping criterion thats based on eigenplanes. This criterion also generates a new candidate if the last wasnt sufficient after finite time. The last criterion is ... [dont remember right now]. In the following section the needed theory from (Guglielmi and Protasov, 2011) is restated.

2.5 Termination conditions

The invariant-polytope algorithm is very efficient but has its caviat. It only is guaranteed to terminate in finite time if the leading eigenvector of the chosen candidate is unique and simple. In the following i will present the according proofs.

2.6 Rebalancing and added starting vertices

Three years after publishing the fundamental algorithm, the creators released a new paper on rebalancing multiple s.m.ps as well as starting with some extra vertices so the polytope is conditioned better (not as flat).

2.7 Eigenvector cases

If the eigenvector is complex a different norm must be used a so called complex symmetric polytope norm. Also in the case of nonnegative matrices a so called cone-norm is used instead of the polytope-norm since its way more efficient.

3 Finite-tree algorithm

In this chapter i introduce the reader to the concepts of the finite-tree algorithm.

3.1 Notation and definitions

Throughout this thesis, we consider a finite set of matrices $\mathcal{A} = \{A_1, \dots, A_J\}$ with $A_j \in \mathbb{R}^{s \times s}$. The computation of the JSR using tree-based methods, requires a structured representation of matrix products.

Encoding of Matrix Products

The set $J^n := \{1, \dots, J\}^n$ denotes the collection of all index sequences of length n . For an index sequence $j = [j_1, \dots, j_n] \in J^n$, the corresponding matrix product is given by

$$A_j = A_{j_n} \cdots A_{j_1}.$$

For $n = 0$, we define the identity matrix $A_j = I$.

Definition of an finite-tree

An (\mathcal{A}, G) -tree is a structured representation of matrix products used to decompose arbitrary matrix products from \mathcal{A} . Lets define it given a set of generator indices $G = \{g_1, \dots, g_I\} \subseteq J^n$:

Definition 3.1.1. *A Tree with the following structure:*

- The root node contains the identity matrix: $t_0 = \{I\}$.
- The root node is parent of exactly J children: $\{A_j P : P \in t\}, j = 1, \dots, J$
- Each node $t \in T$ is either:
 - A leaf (i.e., it has no children),
 - A parent of exactly J children: $\{A_j P : P \in t\}, j = 1, \dots, J$ (positive children),
 - A parent of arbitrarily many generators: $\{A_g^n P : P \in t, n \in \mathbb{N}_0\}$ for some $g \in G$ (negative children).

Is called $(\mathcal{A}, \mathbf{G})$ -tree

Covered Nodes

A node in the tree is called *covered* if it is a subset of one of its ancestors in the tree. Otherwise, it is called *uncovered*. The set of uncovered leaves is denoted as

$$\mathcal{L}(T) := \{L \in t : t \text{ is an uncovered leaf of } T\}.$$

and called leafage of the tree T .

1-Boundedness

An (A, G) -tree T is called *1-bounded* with respect to a sub-multiplicative matrix norm $\|\cdot\|$ if

$$\sup_{L \in \mathcal{L}(T)} \|L\| \leq 1.$$

If a stricter bound holds, i.e., $\sup_{L \in \mathcal{L}(T)} \|L\| < 1$, then the tree is called *strictly 1-bounded*.

Possible trees remark

Since we are only interested in the smallest possible tree (smallest runtime) there are some trees that can be sorted out. For example if we take a generator directly at the beginning of the tree, per definition for the tree to be 1-bounded also choosing $n = 0$, meaning the generator acts as identity, should yield norms less than 1. But that means the generator could have been left out and all norms are still less than 1, making it unnecessary. Also putting a generator as a leaf is inefficient: either the leaf is 1-bounded and the same argument with $n = 0$ from before holds or it is covered by another node. I propose that a generator leaf can only be covered by another generator leaf. That would mean we at least have one 1-bounded generator leaf which is declared inefficient from before.

Proof. Let the encoding of the leaf t_1 node be $j = [P, g]$. P stands for prefix and g is the index for the used generator. For it to be covered, there must exist a node t_2 with the encoding $[P', g, S]$ where the prefix is different and it also is allowed to have a suffix. Since all products from t_1 are in t_2 the suffix can only be of the form $S = g^k$. Now choosing $n = 0$ for the generator of the node t_1 , we get $P = g^k \cdot g^m \cdot P'$ for some $m \in \mathbb{N}$ but that would mean the node at P' has positive and negative children, which is not allowed, forming a contradiction. \square

3.2 Structure of the finite-tree algorithm

We start with the root node, the identity $t_0 = I$. Now we build up a tree that's of the form of 2. Possible methods would be breadth-first- or depth-first-search for building up the tree. The chosen norm in this case is arbitrary but changes the runtime dramatically. If the 1-bounded tree was found, we have proven that there exists a decomposition for every product of matrices from \mathcal{A} .

Algorithm 2 Finite-tree algorithm

Theorem 3.2.1. *If the finite-tree algorithm 2 terminates the $JSR(\mathcal{A})$ was found.*

Proof. The leafage of the generated tree is 1-bounded. Now we take an arbitrary product $P \in \mathcal{A}^n$. If the product is outside the current tree i.e. after applying lexicographic ordering on the encodings of the nodes the encoding of the product is bigger than every other encoding in the tree. That means the structure of (A, G) -trees allows us to find one leaf-node that is a valid prefix of the product. Now we have $P = P'L$ with $\|L\| \leq 1$. Since every leaf-node is guaranteed to have at least length 1 (first branches must be \mathcal{A}) the length of P' is strictly smaller than the length of P . We can repeat that process until $P'^{(k)}$ is within the tree. If the product is within the tree, we can generate a polynomial $p(k)$ that is monotone and an upper bound on the norms of products within the tree, where k is the length of the product. For that we have to analyse the Jordan-Normalform of the products. Since the spectral radius of every matrix in \mathcal{A} and of every generator is less than 1, we now the growth of potencies of the Jordan-blocks is bounded by a polynomial in the order of the size of the Jordan-block. Just taking a maximum over all factors of all such polynomials we generate a polynomial that bounds every product within the tree according to the length of the product. So every product $P \in \mathcal{A}^n$ is of the form $P = S \cdot P_k \cdot \dots \cdot P_1$ where $\|P_i\| \leq 1$ and $\|S\| \leq p(k)$. We imply:

$$JSR(\mathcal{A}) = \lim_k \max_{P \in \mathcal{A}^k} \|P\|^{\frac{1}{k}} \leq \lim_k p(k)^{\frac{1}{k}} = 1$$

The detailed proof can be read in (Möller and Reif, 2014). □

3.3 Efficiency results

Theorem 3.3.1. *If the invariant-polytope algorithm terminates so does the finite-tree algorithm.*

Proof. The proof can be seen in (Möller and Reif, 2014). □

Theorem 3.3.2. *The solution space of the finite-tree algorithm is strictly bigger than the one from the invariant-polytope algorithm.*

Proof. From theorem 3.3.1 we know that the finite-tree algorithm always terminates if the invariant-polytope algorithm does. Lets view an example for when the invariant-polytope algorithm can not terminate due to a missing spectral-gap. But the finite-tree algorithm can find $JSR(\mathcal{A})$ nonetheless. □

Tree generation

The runtime does not only depend on the chosen norm but also on the use of the generators. The idea is that the generators come closest to the maximal growth rate so using them will max-

imize the norm and the subsequent branches can almost only decrease the norms eventually so that the 1-boundedness can be checked earlier.

4 Hybrid approach

In this chapter we want to explore some possible combinations of the before mentioned algorithm schemes and then present the main result of this work, the Tree-flavored-invariant-polytope-algorithm and its termination results.

4.1 Goal and options

In its heart the invariant-polytope algorithm tries to find a norm that is specifically optimized on the given problem, whilst the finite-tree algorithm connects growth to decompositions of products. A clear combination scheme arises naturally, where we use the optimized polytope norm to estimate the products of the finite-tree. From there we can choose a specific order or level of concurrency.

The most modular approach would be to first run the invariant-polytope algorithm for a couple of runs and then use the calculated norm that is specially optimized for the finite-tree algorithm. But that seems to be wasteful since valuable matrix calculations from the finite-tree algorithm could have been used for an even more optimized norm and some polytopes might have already cleared insight for the decompositions that the finite-tree algorithm tries to find. So after a bit of rethinking we managed to come up with a more concurrent algorithm that builds up norms and decompositions in every step.

4.2 Structure of the hybrid algorithm

We try to decompose arbitrary products P from factors from \mathcal{A} , such that their polytope-norms are less than $p(k)$ where k is the number of factors from P and p is a monotone polynomial. This removes the invariance property of the polytope to be built up, since the norms don't have to be less than 1 but it still proves the JSR identity because we take the weighted norms in the length of the products in the three-member-inequality 1.3.

Starting the loop of the invariant-polytope algorithm with a cycle on top that is connected via the generators factors and also the first branches represented by images from the missing $J-1$ factors from \mathcal{A} . Instead of only adding images under vertices from V and matrices from \mathcal{A} directly, from now on we try to find an $(\mathcal{A}, \mathbf{G})$ -tree which is one-bounded i.e. its leafage-polytope-norm is less than 1, for every $v \in V$. For that we generate $(\mathcal{A}, \mathbf{G})$ -tree patterns in the beginning and just go through every remaining vertex and calculate the leafage-norms. From the structure of those trees we can assume that every matrix in \mathcal{A} represents a node for the first branches. For the branches that lead to a leafage-norm less than or equal 1 we are done, for the other branches we have the choice to go deeper or just add some points to V that changes

the leafage-norm of those branches to less than 1. Here we decided to add the points since going deeper just would mean to consider possibly the same products but the tree generation would be more complex with options for depth-first- or breadth-first-search and even using some s.m.p and generator trickery. [might change it in the future]

First points that come to mind are the leafage-points itself since this is what we have tested but generators could be involved meaning there are possibly infinitely many leafage-points. So the next best thing would be the roots of the branches which are guaranteed to be a single matrix from \mathcal{A} . This makes tree generation easy and adds points with likely more distance to the faces of the polytope and makes the norm stronger more quickly.

So in principle for every $v \in V_{\text{rem}}$ take a tree from the generating pool, check the leafage-norm for every root branch, if it is larger than 1 add the point from the root branch to V_{new} and V . After one step change V_{rem} to V_{new} and V_{new} to the empty set and repeat this as long as new vertices have been added. We use V for the polytope-norms and since new points are only being added the norms decrease over time so all 1-bounded trees stay bounded.

After termination the set of trees generated promise a valid decomposition for every product from \mathcal{A} into chunks of norm lesser 1 and one suffix thats of norm less than $p(k)$ for some monotone polynomial, which proofes the question if the chosen radius is maximal.

Algorithm 3 Tree-flavored-invariant-polytope-algorithm

```

 $V := \{v_1, \dots, v_M\}$ 
 $V_{\text{new}} \leftarrow V$ 
while  $V_{\text{new}} \neq \emptyset$  do
     $V_{\text{rem}} \leftarrow V_{\text{new}}$ 
     $V_{\text{new}} \leftarrow \emptyset$ 
    for  $v \in V_{\text{rem}}$  do
        Construct some  $(\mathcal{A}, \mathbf{G})$ -tree  $\mathbf{T}$ 
        for  $L = L'A \in \mathcal{L}(T)$  with  $A \in \mathcal{A}$  do
            if  $\|Lv\|_{\text{cos}(V)} \geq 1$  then
                 $V \leftarrow V \cup Av$ 
                 $V_{\text{new}} \leftarrow V_{\text{new}} \cup Av$ 
return

```

The following figure 4.1 shows an example of the generated tree structure as well as the tested $(\mathcal{A}, \mathbf{G})$ -trees that have been used to bound the products. Here \mathcal{A} consists of two matrices A and B and the chosen candidate is $\Pi = ABA$. At the top you can see the cycle generated by the candidate and the starting vector v_1 as well as branches coming of that are the products that stay in contrast to the cycle. This is what i call the crown and it is generated for every matrix set at the beginning, so no testing has been made so far. All the nodes that are connected through a solid arrow have been added to the vertices V .




Figure 4.1: Cyclic tree structure generated by the algorithm. Vertices added to V (solid arrows) and finite-trees for bounding products (dashed arrows).

Now the finite-tree theory comes in and every vertex gets tested by a $(\mathcal{A}, \mathbf{G})$ -tree (dashed arrows). Take the vector v_6 for example it has been tested by an tree and the branch starting with the matrix B was sufficient i.e. every leaf-node has a polytope-norm of less than 1 (at the time the polytope consists of the vertices $V = \{v_1, \dots, v_6\}$). The branch starting with the matrix A on the other hand didnt, so the vertex $v_7 = Av_6$ has been added to V and the next tree was tested, which was fully sufficient (on every branch). Since the other branches also terminated earlier the algorithm stops, no more vertices are beeing added and the candidate Π was proven to be a s.m.p product.

4.3 Termination results

Theorem 4.3.1. *If Algorithm 3 terminates then $JSR(\mathcal{A}) \leq 1$*

Proof. Suppose the algorithm terminates and creates a set V of vertices then for each $v \in V$ there exists an $(\mathcal{A}, \mathbf{G})$ -Tree \mathbf{T}_v such that $\|Lv\|_{\mathbf{co}_s(V)} \leq 1 \quad \forall L \in \mathcal{L}(\mathbf{T}_v)$.

Taking arbitrary product $P \in \mathcal{A}^k$ and vector $v \in \mathbf{co}_s(V)$ we get:

$$\|Pv\| = \|P \sum \lambda_i v_i\| \leq \sum |\lambda_i| \|Pv_i\| \quad \text{with} \quad \sum |\lambda_i| \leq 1$$

Now for every $i = 1 \dots |V|$ there either exists a Tree-partition where $P = P'_i L_i$ with $L_i \in \mathcal{L}(T_{v_i})$ and $\|L_i v_i\| \leq 1$ or P is element of a node of T_{v_i} thats not a leaf. If the partition exists

the vector $L_i v_i$ lands within the symmetrized convex hull of V and thus has a linear combination like before.

$$\|Pv_i\| = \|P'_i(L_i v_i)\| \leq \sum |\mu_j| \|P'_i v_j\| \quad \text{with} \quad \sum |\mu_j| \leq 1$$

If there exists no such partition then the product is in some sense already small enough: A similar argument, like in (Möller and Reif, 2014) can be made, that theres only finitely many nodes in every tree and the spectral radius of every factor of such a matrix-product is less then 1. Norms of such products can be bounded by a monotone increasing polynomial $p(k)$ in the amount of factors k .

Since every product can be reduced to a product that has less factors (trees have to have at least \mathcal{A} as branches) until it lies within an according tree we eventually get:

$$\begin{aligned} \|Pv\| &\leq \sum_i |\lambda_i| \|Pv_i\| \leq \sum_{i_1} \sum_{i_2} |\lambda_{i_1}| |\lambda_{i_2}| \|P'_{i_1} v_{i_2}\| \leq \dots \\ &\leq \sum_{i_1, \dots, i_k \in \Sigma} |\lambda_{i_1}| \cdot \dots \cdot |\lambda_{i_k}| \|P'_{i_1, \dots, i_k} v_{i_k}\| \leq \sum_{i_1, \dots, i_k \in \Sigma} |\lambda_{i_1}| \cdot \dots \cdot |\lambda_{i_k}| \cdot p(k) \leq p(k) \end{aligned}$$

the last inequality holds since the factors λ_α come from the linear combination of the polytopes vertices and the norm is less than 1 so the factors fullfil $\sum |\lambda_\alpha| \leq 1$. This implies $\|P\| \leq p(k)$ thus, since P was chosen arbitrary:

$$JSR(\mathcal{A}) = \lim_k \max_{P \in \mathcal{A}^k} \|P\|^{\frac{1}{k}} \leq \lim_k p(k)^{\frac{1}{k}} = 1$$

just like with the finite-tree algorithm. □

The idea for the polynomial is that by knowing the spectral radius of every matrix and generator to be less than 1 the jordan-blocks of matrices from one node are predictable in their growth under exponantiation. The growth is less than a monotone polynomial in the order of the size of the jordan-block. Now we have such a polynomial for every node and there are only finitely many nodes so we take the maximum of every monome factor and create a polynomial that bounds every nodes products.

This only works since the number of nodes is finite. If we would apply the same strategy on the infinite tree the absolute values of the factors for the polynomial would not be bounded.

Corollary 4.3.2. *If the chosen trees are always $T_v = \mathcal{A}$ then we would have exactly the invariant-polytope algorithm 1. Which makes the new approach a real generalization.*

Corollary 4.3.3. *If the chosen trees do not make use of generators, then the polynomial bounding would collapse to just a constant $K = \max_{v \in V} \max_{P \in T_v} \|P\|_{\text{cos}(V)}$ that is the maximal polytope-norm over all nodes and all used $(\mathcal{A}, \mathbf{G})$ -trees. This constant would not depend on the sizes of the products anymore.*

4.4 Efficiency

Theorem 4.4.1. *The Tree-flavored-invariant-polytope algorithm has a bigger solution space than the invariant-polytope algorithm.*

empty. □

Theorem 4.4.2. *The Tree-flavored-invariant-polytope algorithm is numerically faster than the finite-tree algorithm starting with an ill-conditioned norm.*

It is hard to proof this but in the next Chapter 5 i will present some numerical results that support the claim.

Remarks

Used trees

It was proven in (Möller and Reif, 2014) that in order to have a bigger solution space then algorithm 1 the use of generators is mandatory. The choice of the testing-trees is not trivial and subject to active research.

Stopping criterions

Some of the stopping criterions of the invariant-polytope algorithm can still be used as the structure is very similar. Just some remarks...

Rebalancing and added starting vertices

Rebalancing allows for multiple s.m.p's in the invariant-polytope algorithm so it has a bigger solution space but it is still comparable to the hybrid approach due to ...

Polytope invariance

Since the norms of matrices in \mathcal{A} are allowed to be bigger than 1 just less than $p(1)$ it is possible the polytope is not invariant. Although this is not the main goal, as we are trying to compute $JSR(\mathcal{A})$, but if one is interested in finding an invariant polytope anyway, one has to ...

5 Numerics

5.1 Runtime of algorithm 3 on generated matrices

5.1.1 Varying matrix dimensions

5.1.2 Varying matrix condition

5.1.3 Varying matrix s.m.p number

5.2 Comparison algorithm 3 and 1

5.2.1 Big matrix dimensions

5.2.2 Multiple s.m.p's

5.2.3 old closed problems couldnt be solved by 1

5.3 Comparison algorithm 3 and 2

5.3.1 Multiple s.m.p's and big matrix dimensions

5.3.2 old closed problems couldnt be solved by 2

5.4 Solving open problems

6 Conclusion

6.1 New approach

6.2 Solved Problems

6.3 Efficiency

6.4 Unsolved problems

6.5 Further research

6.5.1 Complex case

6.5.2 Nonnegative case

6.5.3 Optimizing tree generation and vertex choice

Bibliography

- Blondel, V. D. and Tsitsiklis, J. N. (2000). A survey of computational complexity results in systems and control. *Automatica*, 36(9):1249–1274.
- Gripenberg, G. (1996). Computing the joint spectral radius. *Linear Algebra and its Applications*, 234:43–60.
- Guglielmi, N. and Protasov, V. (2011). Exact computation of joint spectral characteristics of linear operators.
- Guglielmi, N. and Protasov, V. (2013). Exact Computation of Joint Spectral Characteristics of Linear Operators. *Foundations of Computational Mathematics*, 13(1):37–97.
- Guglielmi, N. and Protasov, V. Y. (2015a). Invariant polytopes of linear operators with applications to regularity of wavelets and of subdivisions. *none*.
- Guglielmi, N. and Protasov, V. Y. (2015b). Invariant polytopes of linear operators with applications to regularity of wavelets and of subdivisions.
- Heil, C. (1993). Ten Lectures on Wavelets (Ingrid Daubechies). *SIAM Review*, 35(4):666–669.
- Jungers, R. M. (2009). The Joint Spectral Radius. *none*.
- Mejstrik, T. (2020a). Algorithm 1011: Improved Invariant Polytope Algorithm and Applications. *ACM Transactions on Mathematical Software*, 46(3):1–26.
- Mejstrik, T. (2020b). Improved invariant polytope algorithm and applications.
- Mejstrik, T. and Reif, U. (2024a). A Hybrid Approach to Joint Spectral Radius Computation. *none*.
- Mejstrik, T. and Reif, U. (2024b). Hybrid approach to the joint spectral radius computation.
- Möller, C. and Reif, U. (2014). A tree-based approach to joint spectral radius determination. *Linear Algebra and its Applications*, 463:154–170.
- Rota, G.-C. and Gilbert Strang, W. (1960). A note on the joint spectral radius. *Indagationes Mathematicae (Proceedings)*, 63:379–381.