MASTERARBEIT | MASTER'S THESIS

Titel | Title

Termination results for hybrid approach of Joint Spectral Radius computation

verfasst von | submitted by

Aaron Pumm, B.Sc.

angestrebter akademischer Grad | in partial fulfilment of the requirements for the degree of

Master Zusatz M.Sc.

Wien, | Vienna, 2026

Studienkennzahl lt. Studienblatt | degree programme code as it appears on the student record sheet:	UA 066 821
Studienrichtung lt. Studienblatt | degree programme as it appears on the student record sheet:	Masterstudium Mathematik
Betreut von | Supervisor:	Assoz. Prof. Dr. Martin Ehler
Ko-betreut von | Co-Supervisor:	Dr. Thomas Mejstrik MMag.

Contents

1. Introduction and Background **6**
1.1. Notation and Preliminaries . . . 7
1.2. Motivation and Applications of the Joint Spectral Radius . . . 7
1.3. Fundamental Theoretical Results . . . 9
1.4. Preprocessing and Normalization . . . 15

2. Invariant-Polytope Algorithm **16**
2.1. Overview . . . 16
2.2. Algorithmic Formulation . . . 16
2.3. Termination Theory . . . 18
2.4. Implementation Details . . . 20
2.5. Remarks . . . 21

3. Finite-Tree Algorithm **22**
3.1. Overview . . . 22
3.2. Algorithmic Formulation . . . 24
3.3. Termination Theory . . . 25
3.4. Implementation Details . . . 27
3.5. Remarks . . . 27

4. Invariant-Tree Algorithm **28**
4.1. Overview . . . 28
4.2. Algorithmic Formulation . . . 29
4.3. Termination Theory . . . 29
4.4. Implementation Considerations . . . 33
4.5. Remarks . . . 34

5. Numerics **36**
5.1. Solves on generated Matrix Sets . . . 36
5.2. Solved Problems with One s.m.p. . . . 36
5.3. Solved Problems with Multiple s.m.p.s . . . 36

6. Conclusions and Outlook **37**
6.1. Summary of Invariant-Tree numerics . . . 37
6.2. Resolved Problem Classes . . . 37
6.3. Remaining Open Problem Classes . . . 37

6.4. Directions for Future Research 37

A. Notation Glossary **38**
A.1. Constants and Globally Defined Variables 38
A.2. Products and Multi-Indices 38
A.3. Norms and JSR 39
A.4. Algorithmic Notation 39
A.5. Tree Terminology 40

Zusammenfassung

Diese Arbeit behandelt die Berechnung des gemeinsamen Spektralradius (JSR) einer endlichen Menge von Matrizen. Die Kapitel zum Invariant-Polytope-Algorithmus und zum Finite-Tree-Algorithmus fassen den bestehenden Forschungsstand zusammen und dienen als Grundlage für den hybriden Ansatz. Der originäre Beitrag dieser Arbeit liegt in der detaillierten Ausarbeitung des Invariant-Tree-Algorithmus (Algorithmus 3) sowie seiner Terminierungstheorie auf Basis von (Mejstrik and Reif, 2025). Implementierungsdetails zur endlichen 1-Beschränkung erzeugter Produktmengen werden aus der Literatur übernommen und nicht als neuer Beitrag beansprucht. Zusätzlich werden numerische JSR-Rechnungen für Matrizenmengen gezeigt, bei denen bestehende Methoden scheitern oder ineffizient sind.

Abstract

This thesis studies computation of the joint spectral radius (JSR) for finite matrix families. Chapters on the invariant-polytope algorithm and the finite-tree algorithm are literature-based summaries that provide the theoretical background for the hybrid method. The original contribution of this thesis is the detailed formulation of the invariant-tree algorithm (Algorithm 3) and its termination theory, based on (Mejstrik and Reif, 2025). Implementation details for finite 1-bounding of generated product sets are adopted from the cited literature and are not claimed as original here. The thesis also includes numerical experiments on matrix families where standalone methods are ineffective or computationally expensive.

Keywords: Joint spectral radius, invariant-tree algorithm, invariant-polytope algorithm, finite-tree algorithm, matrix norms.

1. Introduction and Background

The *Joint Spectral Radius* (JSR) was first introduced by G.-C. Rota and G. Strang in 1960 (Rota and Strang, 1960). They described the JSR as the maximal exponential growth rate of a product of matrices taken from a finite set. Since its inception, the JSR has become a cornerstone in various mathematical and engineering disciplines due to its ability to encapsulate the asymptotic behavior of matrix products.

The concept gained significant traction in the 1990s when researchers began exploring its theoretical properties and practical implications. Notable advancements include its application in wavelet theory, where it assists in the construction of refinable functions (Daubechies and Lagarias, 1992), as well as in control theory, where it is used to analyze the stability of switched linear systems (Jungers, 2009), for which we provide a small example below. The computational challenges associated with determining the JSR have inspired the development of several algorithms, such as the invariant-polytope algorithm (Guglielmi and Protasov, 2013) and the finite-tree algorithm (Möller and Reif, 2014).

Despite the progress, JSR computation remains challenging, particularly due to the exponential complexity of exploring all possible matrix products.

This thesis has a two-part scope. Chapters 2 and 3 summarize established theory and algorithms from the literature and are included as background for the hybrid framework. The original contribution of this thesis is the detailed formulation of Algorithm 3 (invariant-tree algorithm) and its termination theory in Chapter 4, developed on top of (Mejstrik and Reif, 2025). Implementation details for finite 1-bounding of generated product sets are adopted from existing work and are not claimed as original here.

Thesis Outline

This thesis is structured as follows: Chapter 1 provides background on the JSR and its basic properties. Chapters 2 and 3 provide literature summaries of the invariant-polytope and finite-tree approaches, including the required theoretical foundations. Chapter 4 contains the thesis-specific contribution: a detailed invariant-tree algorithmic formulation and corresponding termination theory. In related literature, this construction is also called the *tree-flavoured-invariant-polytope-algorithm* (Mejstrik and Reif, 2025). Chapter 5 presents numerical results to analyze the efficiency and applicability of the invariant-tree algorithm. Chapter 6 concludes with insights and future directions. Appendix A collects a complete notation glossary for quick reference.

1.1. Notation and Preliminaries

Unless stated otherwise, all matrices are real $d \times d$ matrices and all matrix norms are submultiplicative. We use the following notation conventions consistently across all chapters:

- Matrix families and dimensions:

$$\mathcal{A} = \{A_1, \ldots, A_n\}, \qquad \tilde{\mathcal{A}} = \{\tilde{A}_1, \ldots, \tilde{A}_n\}, \qquad d = \text{matrix dimension}, \quad n = |\mathcal{A}|.$$

- Generator notation:

$$\mathcal{G} = \{G_1, \ldots, G_m\}, \qquad m = |\mathcal{G}|.$$

- Positive indices and products:

$$\mathcal{I} := \{1, \ldots, n\}, \quad I = [i_1, \ldots, i_k] \in \mathcal{I}^k, \quad A_I := A_{i_k} \cdots A_{i_1}, \quad \mathcal{A}^k := \{A_I : I \in \mathcal{I}^k\}.$$

- Scalars and spectra: α denotes a scalar, x a vector, λ an eigenvalue, $\rho(\cdot)$ the spectral radius, and $\sigma(\cdot)$ a singular value.
- Polytope norms:

$$\|\cdot\|_{\mathrm{co_s}(V)} = \|\cdot\|_V.$$

Appendix A contains the full notation glossary used throughout the thesis.

1.2. Motivation and Applications of the Joint Spectral Radius

Let $\mathcal{A} = \{A_1, \cdots, A_n\}$ be a finite set of $d \times d$ matrices. Consider a discrete-time linear system generated by a matrix sequence from $\mathcal{A}$:

$$x_{t+1} = B_t x_t, \quad x_0 \in \mathbb{R}^d, \tag{1.1}$$

where $B_t \in \mathcal{A}$ for all $t \in \mathbb{N}$.

A key question is whether the system is *uniformly asymptotically stable*, i.e., whether every trajectory (x_t) converges to zero for all initial states and all matrix sequences $(B_t)_{t\geq 0}$ with $B_t \in \mathcal{A}$:

$$\lim_{t\to\infty} x_t = 0.$$

This stability question is answered by a generalization of the spectral radius for a single matrix, the so-called *joint spectral radius (JSR)* of a matrix set $\mathcal{A}$ (Rota and Strang, 1960).

Definition 1.2.1. (Rota and Strang, 1960) For a finite set of matrices $\mathcal{A} = \{A_1, \ldots, A_n\}$, the JSR is defined as:

$$\mathrm{JSR}(\mathcal{A}) := \lim_{k\to\infty} \max_{P\in\mathcal{A}^k} \|P\|^{1/k}. \tag{1.2}$$

where $\|\cdot\|$ denotes any matrix norm.

This definition quantifies the maximal asymptotic growth rate of matrix products from $\mathcal{A}$.

Illustrative Example

Consider the matrix set:

$$\mathcal{A} = \left\{ A_1 = \begin{bmatrix} 0.5 & 1 \\ 0 & 0.5 \end{bmatrix}, \quad A_2 = \begin{bmatrix} 0.3 & 0 \\ 1 & 0.3 \end{bmatrix} \right\}.$$

Each matrix individually has spectral radius less than 1. However, their product

$$A_1 A_2 = \begin{bmatrix} 0.5 & 1 \\ 0 & 0.5 \end{bmatrix} \begin{bmatrix} 0.3 & 0 \\ 1 & 0.3 \end{bmatrix} = \begin{bmatrix} 1.15 & 0.3 \\ 0.5 & 0.15 \end{bmatrix}$$

has spectral radius of approximately

$$\rho(A_1 A_2) \approx 1.2825 > 1.$$

This example shows that switching can amplify the effects of individual matrices, leading to divergent behavior even if all single matrices are contractive. Thus, the JSR can exceed the spectral radius of any individual matrix in the set.

The following results motivate the importance of bounding or approximating $\mathrm{JSR}(\mathcal{A})$ (Jungers, 2009):

- The system (1.1) is *uniformly bounded* (i.e., $\sup_t \|x_t\| < \infty$ for all $x_0 \in \mathbb{R}^d$) if and only if $\mathrm{JSR}(\mathcal{A}) \leq 1$ and $\mathcal{A}$ is irreducible.
- It is *uniformly asymptotically stable* if and only if $\mathrm{JSR}(\mathcal{A}) < 1$.

Understanding the joint spectral radius is therefore critical for analyzing the asymptotic behavior of switched linear systems and motivates reliable computational methods (Jungers, 2009). Beyond control theory, the JSR also appears in approximation theory and subdivision theory, with foundational references including stationary subdivision and two-scale refinement equations (Cavaretta et al., 1991; Daubechies and Lagarias, 1992; Maesumi, 2000). Related spectral-growth ideas are used in coding-theoretic capacity problems, including codes that avoid specified differences (Moision et al., 2001; Blondel et al., 2006), and in discrete mathematics and combinatorics (Protasov, 2000;

Jungers et al., 2008, 2009). These links motivate efficient computational methods, including the invariant-tree algorithm studied in this thesis (Mejstrik and Reif, 2025).

1.3. Fundamental Theoretical Results

Lemma 1.3.1 (Fekete's Lemma)**.** *Let $(a_n)_{n\geq 1}$ be a sequence of real, nonnegative numbers such that*

$$a_{m+n} \leq a_m + a_n \qquad \text{for all } m, n \geq 1. \tag{1.3}$$

Then the limit

$$\lim_{n\to\infty} \frac{a_n}{n}$$

exists and is equal to $\inf_{n\geq 1}\left(\frac{a_n}{n}\right)$.

Proof. Set $b_n := a_n/n$ and let $\alpha := \inf_{n\geq 1} b_n$. Fix $n \geq 1$ and write any $k \geq 1$ as $k = qn+r$ with integers $q \geq 0$ and $0 \leq r < n$. By subadditivity (1.3),

$$a_k = a_{qn+r} \leq qa_n + a_r.$$

Dividing by k gives

$$\frac{a_k}{k} \leq \frac{qa_n}{qn+r} + \frac{a_r}{qn+r} \leq \frac{a_n}{n} + \frac{a_r}{qn}.$$

Taking $\limsup_{k\to\infty}$ (equivalently $q \to \infty$) yields

$$\limsup_{k\to\infty} \frac{a_k}{k} \leq \frac{a_n}{n}.$$

Since n was arbitrary, we obtain

$$\limsup_{k\to\infty} \frac{a_k}{k} \leq \inf_{n\geq 1} \frac{a_n}{n} = \alpha.$$

On the other hand, by definition of infimum we always have $\alpha \leq a_k/k$ for all k, hence $\alpha \leq \liminf_{k\to\infty} a_k/k$. Therefore $\liminf = \limsup = \alpha$, so the limit exists and equals α. □

Theorem 1.3.2. *(Rota and Strang, 1960; Jungers, 2009) The JSR is well-defined and independent of the choice of the matrix norm.*

Proof. Let $\|\cdot\|_1$ and $\|\cdot\|_2$ be two matrix norms on $\mathbb{R}^{n\times n}$. By equivalence of norms in finite-dimensional vector spaces, there exist constants $c, C > 0$ such that:

$$c\|P\|_1 \leq \|P\|_2 \leq C\|P\|_1 \quad \text{for all } P \in \mathbb{R}^{d\times d}.$$

Therefore, we have

$$\begin{aligned}
&\lim_{k\to\infty} \max_{P\in\mathcal{A}^k} \|P\|_1^{1/k} \leq \lim_{k\to\infty} \left(\frac{1}{c}\right)^{\frac{1}{k}} \max_{P\in\mathcal{A}^k} \|P\|_2^{1/k} \\
&= \lim_{k\to\infty} \max_{P\in\mathcal{A}^k} \|P\|_2^{1/k} \leq \lim_{k\to\infty} C^{\frac{1}{k}} \max_{P\in\mathcal{A}^k} \|P\|_1^{1/k} \\
&= \lim_{k\to\infty} \max_{P\in\mathcal{A}^k} \|P\|_1^{1/k},
\end{aligned}$$

which makes the limit independent of the choice of the used matrix norm. The existence of this limit is a direct consequence of Fekete's Lemma. □

To facilitate the analysis of the algorithms and presentation of the main results, we first establish fundamental properties of the JSR.

Homogeneity of the Joint Spectral Radius

Proposition 1.3.3. *(Jungers, 2009) The JSR is homogeneous, meaning for any set of matrices $\mathcal{A}$ and scalar $\alpha \in \mathbb{R}$ we have*

$$\mathrm{JSR}(\alpha\mathcal{A}) = |\alpha| \,\mathrm{JSR}(\mathcal{A}). \tag{1.4}$$

Proof. Let $\mathcal{A} = \{A_1, \ldots, A_n\}$ and $\alpha \in \mathbb{R}$. Then:

$$\begin{aligned}
\mathrm{JSR}(\alpha\mathcal{A}) &= \lim_{k\to\infty} \max_{A_i\in\mathcal{A}^k} \|\alpha A_{i_k} \cdots \alpha A_{i_1}\|^{1/k} \\
&= |\alpha| \lim_{k\to\infty} \max_{A_i\in\mathcal{A}^k} \|A_{i_k} \cdots A_{i_1}\|^{1/k} \\
&= |\alpha| \,\mathrm{JSR}(\mathcal{A})
\end{aligned}$$

□

Berger–Wang Characterization and Three-Member Inequality

Proposition 1.3.4. *(Berger and Wang, 1992) The* JSR *can equivalently be defined as*

$$\mathrm{JSR}(\mathcal{A}) = \limsup_{k\to\infty} \max_{P\in\mathcal{A}^k} \rho(P)^{1/k},$$

if $\mathcal{A}$ is a finite (or compact) set.

Proposition 1.3.5. *(Rota and Strang, 1960; Jungers, 2009) Let $\|\cdot\|$ be a matrix norm. Then for every $k \in \mathbb{N}$,*

$$\max_{P\in\mathcal{A}^k} \rho(P)^{1/k} \;\leq\; \mathrm{JSR}(\mathcal{A}) \;\leq\; \max_{P\in\mathcal{A}^k} \|P\|^{1/k}. \tag{1.5}$$

Proof. Fix $k \in \mathbb{N}$ and choose

$$P_{\max} \in \arg\max_{P \in \mathcal{A}^k} \rho(P)^{1/k}.$$

For the left-hand side we obtain, for any $r \in \mathbb{N}$,

$$\begin{aligned} \max_{P \in \mathcal{A}^k} \rho(P)^{1/k} &= \rho(P_{\max})^{1/k} = \rho(P_{\max})^{r/(kr)} \\ &= \rho(P_{\max}^r)^{1/(kr)} \leq \max_{P \in \mathcal{A}^{kr}} \rho(P)^{1/(kr)} \\ &\leq \limsup_{r \to \infty} \max_{P \in \mathcal{A}^{kr}} \rho(P)^{1/(kr)} \leq \mathrm{JSR}(\mathcal{A}). \end{aligned} \tag{1.6}$$

For the right-hand side, set

$$M_n := \max_{P \in \mathcal{A}^n} \|P\| \qquad (n \in \mathbb{N}).$$

By submultiplicativity of the norm, every $Q \in \mathcal{A}^{m+n}$ can be written as $Q = RS$ with $R \in \mathcal{A}^m$ and $S \in \mathcal{A}^n$, hence

$$\|Q\| \leq \|R\| \, \|S\| \leq M_m M_n, \tag{1.7}$$

and taking maxima over $Q \in \mathcal{A}^{m+n}$ yields

$$M_{m+n} \leq M_m M_n \qquad \text{for all } m, n \in \mathbb{N}. \tag{1.8}$$

Define the real sequence $a_n := \log M_n$. Then from (1.8),

$$a_{m+n} = \log M_{m+n} \leq \log(M_m M_n) = \log M_m + \log M_n = a_m + a_n, \tag{1.9}$$

so $(a_n)_{n \geq 1}$ is subadditive. By Fekete's lemma, the limit $\lim_{n \to \infty} a_n/n$ exists and equals $\inf_{n \geq 1} a_n/n$. In particular, for the fixed k we have

$$\lim_{n \to \infty} \frac{a_n}{n} = \inf_{n \geq 1} \frac{a_n}{n} \leq \frac{a_k}{k}. \tag{1.10}$$

Exponentiating (1.10) gives

$$\begin{aligned} \lim_{n \to \infty} M_n^{1/n} &= \exp\left(\lim_{n \to \infty} \frac{a_n}{n} \right) \leq \exp\left(\frac{a_k}{k} \right) \\ &= \exp\left(\frac{\log M_k}{k} \right) = M_k^{1/k} \\ &= M_k^{1/k} = \max_{P \in \mathcal{A}^k} \|P\|^{1/k}. \end{aligned} \tag{1.11}$$

Since $\mathrm{JSR}(\mathcal{A}) = \lim_{n \to \infty} \max_{P \in \mathcal{A}^n} \|P\|^{1/n} = \lim_{n \to \infty} M_n^{1/n}$ by definition, the right-hand

inequality in 1.5 follows. □

This inequality underlies many computational methods, since it yields computable lower and upper bounds on $\mathrm{JSR}(\mathcal{A})$. As $k \to \infty$, the upper bound $\max_{P \in \mathcal{A}^k} \|P\|^{1/k}$ converges to $\mathrm{JSR}(\mathcal{A})$. The lower bounds $\max_{P \in \mathcal{A}^k} \rho(P)^{1/k}$ need not converge, but their limit superior equals $\mathrm{JSR}(\mathcal{A})$ by Proposition 1.3.4.

Similarity Invariance and Reducibility

Proposition 1.3.6. *(Jungers, 2009) The* JSR *is invariant under similarity transformations, i.e., for any invertible matrix T we have*

$$\mathrm{JSR}(\mathcal{A}) = \mathrm{JSR}(T^{-1}\mathcal{A}T). \tag{1.12}$$

Proof. This follows from the definition of the JSR and the properties of the spectral radius.

$$\begin{aligned}
\mathrm{JSR}(T^{-1}\mathcal{A}T) &= \limsup_{k\to\infty} \max_{A_i \in \mathcal{A}^k} \rho(T^{-1}A_{i_k}T \cdot \ldots \cdot T^{-1}A_{i_1}T)^{1/k} \\
&= \limsup_{k\to\infty} \max_{A_i \in \mathcal{A}^k} \rho(T^{-1}A_{i_k} \cdot \ldots \cdot A_{i_1}T)^{1/k} \\
&= \limsup_{k\to\infty} \max_{A_i \in \mathcal{A}^k} \rho(A_{i_k} \cdot \ldots \cdot A_{i_1})^{1/k} \\
&= \mathrm{JSR}(\mathcal{A})
\end{aligned}$$

□

Definition 1.3.7. (Jungers, 2009) A set of matrices is called *(commonly) reducible* if there exists a nontrivial subspace of $\mathbb{R}^n$ that is invariant under all matrices in the set. This means there exists a change of basis that block-triangularizes all matrices in $\mathcal{A}$ at the same time. If $\mathcal{A}$ is not reducible, it is called *irreducible*.

Proposition 1.3.8. *(Protasov, 1997) If $\mathcal{A}$ is reducible, then*

$$\mathrm{JSR}(\mathcal{A}) = \max\{\mathrm{JSR}(\mathcal{B}), \mathrm{JSR}(\mathcal{D})\}, \tag{1.13}$$

where $\mathcal{B} = \{B_1, \ldots, B_n\}$ and $\mathcal{D} = \{D_1, \ldots, D_n\}$ are the blocks of the block-triangularized matrices

$$T^{-1}A_jT = \begin{bmatrix} B_j & C_j \\ 0 & D_j \end{bmatrix}$$

under some invertible matrix T.

Example 1.3.9.

$$A_1 = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}, \quad A_2 = \begin{bmatrix} 3 & 2 \\ 2 & 3 \end{bmatrix}$$

These matrices are reducible because there exists a nontrivial change of basis that transforms them into block-triangular form. The eigenvectors common to both matrices are

$$v_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix}, \quad v_2 = \begin{bmatrix} 1 \\ -1 \end{bmatrix}.$$

Using the change of basis matrix

$$T = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix},$$

we compute the similarity transformations

$$T^{-1}A_1T = \begin{bmatrix} 3 & 0 \\ 0 & 1 \end{bmatrix}, \quad T^{-1}A_2T = \begin{bmatrix} 5 & 0 \\ 0 & 1 \end{bmatrix}.$$

Since both matrices are upper triangular in this basis, they share a common invariant subspace, proving that they are reducible. Now, from Proposition 1.3.8, we know that the JSR can be calculated as

$$\mathrm{JSR}(\mathcal{A}) = \max\{\mathrm{JSR}(\{3,5\}), \mathrm{JSR}(\{1,1\})\} = 5$$

This can be applied iteratively until the sets of blocks are all irreducible. The problem is thereby split into similar problems of smaller dimension. For the following considerations we can now assume $\mathcal{A}$ to be irreducible.

Finiteness Property and Spectrum-Maximizing Products

Definition 1.3.10. (Lagarias and Wang, 1995) A matrix set $\mathcal{A} = \{A_1, \ldots, A_n\}$ is said to possess the *finiteness property* if there exists a finite sequence of matrices $A_{i_1}, \ldots, A_{i_k}$ such that

$$\mathrm{JSR}(\mathcal{A}) = \|A_{i_k} \cdots A_{i_1}\|^{1/k}. \tag{1.14}$$

Any such product is referred to as a *spectrum-maximizing product* (s.m.p.).

Although the finiteness property does not hold for all matrix sets, it is crucial for the effectiveness and termination of most algorithmic approaches. This is shown, for example, in (Blondel et al., 2003).

Computational Complexity of JSR Computation

The computation of the joint spectral radius (JSR) is theoretically challenging due to several complexity results.

Approximating the JSR is NP-hard (Tsitsiklis and Blondel, 1997), meaning no polynomial-time algorithm is expected unless P = NP. Moreover, for general sets of matrices, even deciding whether the JSR is strictly less than one is undecidable (Blondel and Tsitsiklis, 2000). Despite these limitations, many structured instances still allow efficient numerical approximation.

S.m.p.-Candidates and Generators

Approximating the JSR, using the invariant-polytope or the finite-tree algorithms, requires identifying candidate products or *generators* of the matrix set that contribute most significantly to the asymptotic growth rate. These generators are often derived through optimization techniques and their identification is a key step in computational algorithms.

Spectral Gap

Definition 1.3.11. (Mejstrik and Reif, 2025, Definition 1.9) A finite matrix family $\mathcal{A}$ is said to possess a *spectral gap* if there exists a constant $\gamma \in [0, 1)$ such that for every $k \in \mathbb{N}$ and every product $A_I \in \mathcal{A}^k$,

$$\rho(A_I) = \mathrm{JSR}(\mathcal{A})^k \quad \text{or} \quad \rho(A_I) \leq \gamma \ \mathrm{JSR}(\mathcal{A})^k.$$

Equivalently, every non-s.m.p. product is uniformly separated from the maximal growth rate by the factor $\gamma < 1$.

Remark 1.3.12. (Mejstrik and Reif, 2025; Guglielmi and Protasov, 2024) In the invariant-polytope setting, a spectral gap is a key finite-termination assumption together with uniqueness (up to cyclic permutations and powers) and simplicity of the leading eigenvalue of the s.m.p. candidate. Later chapters show where this gap condition enters the algorithmic termination analysis. We also note that (Guglielmi and Protasov, 2024) studies a related notion of gap as separation between the first and second Lyapunov exponents for matrix families.

Example 1.3.13. Consider

$$A_1 = \begin{bmatrix} 2 & 0 \\ 0 & 1 \end{bmatrix}, \qquad A_2 = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}, \qquad \mathcal{A} = \{A_1, A_2\}.$$

Since $\max\{\rho(A_1), \rho(A_2)\} = 2$ and $\|A_1\|_\infty = \|A_2\|_\infty = 2$, we get $\mathrm{JSR}(\mathcal{A}) = 2$. Any product $A_I \in \mathcal{A}^k$ has the form

$$A_I = A_1^m A_2^{k-m} = \begin{bmatrix} 2^m & 0 \\ 0 & 2^{k-m} \end{bmatrix},$$

so

$$\rho(A_I) = 2^{\max\{m, k-m\}}.$$

If $m = 0$ or $m = k$, then $\rho(A_I) = 2^k = \mathrm{JSR}(\mathcal{A})^k$ and A_I is an s.m.p. product. If $1 \leq m \leq k-1$, then

$$\rho(A_I) \leq 2^{k-1} = \tfrac{1}{2}\ \mathrm{JSR}(\mathcal{A})^k.$$

Hence Definition 1.3.11 holds with $\gamma = \frac{1}{2}$.

1.4. Preprocessing and Normalization

Both algorithms use the same preprocessing pipeline:

We want to find the JSR of the finite set of matrices $\mathcal{A} = \{A_1, \cdots, A_n\}$.

1. **Assumptions**: $\mathcal{A}$ is irreducible.
2. **Candidates**: Efficiently find products $P = A_{i_k} \cdots A_{i_1}$ of matrices from $\mathcal{A}$ that maximize the averaged spectral radius $\lambda := \rho(P)^{\frac{1}{k}}$.
3. **Rescaling**: Transform $\mathcal{A} \to \tilde{\mathcal{A}}$ with $\tilde{A}_i := \frac{1}{\lambda} A_i$.
4. **Verification**: Establish that $\mathrm{JSR}(\tilde{\mathcal{A}}) = 1$ using the three-member inequality. By homogeneity, this is equivalent to $\mathrm{JSR}(\mathcal{A}) = \lambda$.

The considered algorithms differ only in step 4. The invariant-polytope algorithm tries to find a vector norm whose induced matrix norm already bounds the matrices from $\mathcal{A}$ sufficiently. The finite-tree algorithm, on the other hand, bounds products using a partitioning space that separates each product into terms bounded by 1 and a remainder term that does not grow fast enough to overcome the k-th root in the JSR definition (polynomial growth).

2. Invariant-Polytope Algorithm

This chapter is a literature-based summary of the invariant-polytope approach and related termination theory. It is included as background for Chapter 4 and does not claim original algorithmic or theoretical contributions.

2.1. Overview

2.1.1. Motivating Ideas

The invariant-polytope algorithm aims to certify an upper bound for the joint spectral radius by constructing a norm that is adapted to the matrix family.

2.1.2. Problem Formulation

After preprocessing (Section 1.4), we have a scaled family $\tilde{\mathcal{A}}$ and a candidate product for the spectral maximizing product (s.m.p.). The remaining task is to certify $\mathrm{JSR}(\tilde{\mathcal{A}}) \leq 1$, which implies that the candidate scaling is correct.

2.1.3. Core Idea

The core idea is to construct a balanced polytope $\mathcal{P} \subset \mathbb{R}^d$ that is invariant under $\tilde{\mathcal{A}}$. Its Minkowski functional then yields an extremal norm for $\mathcal{A}$, and this is sufficient to prove that the chosen candidate is indeed an s.m.p.

2.2. Algorithmic Formulation

2.2.1. Algorithm Description

Starting from an initial finite set of vertices

$$V_0 = \{v_1, \ldots, v_M\},$$

the algorithm iteratively enlarges V by testing images Av with $A \in \tilde{\mathcal{A}}$ and current vertices v. If $\|Av\|_{\mathrm{co_s}(V)} \geq 1$, the point Av is appended. If one full iteration yields no new vertices, the procedure stops and returns $\mathrm{co_s}(V)$.

2.2.2. Pseudocode

Algorithm 1 Invariant-polytope algorithm

```
V := {v_1, ··· , v_M}
V_new ← V
while V_new ≠ ∅ do
    V_rem ← V_new
    V_new ← ∅
    for v ∈ V_rem do
        for A ∈ 𝒜 do
            if ‖Av‖_{co_s(V)} ≥ 1 then
                V ←∪ Av
                V_new ←∪ Av
return co_s(V)
```

2.2.3. Visualization

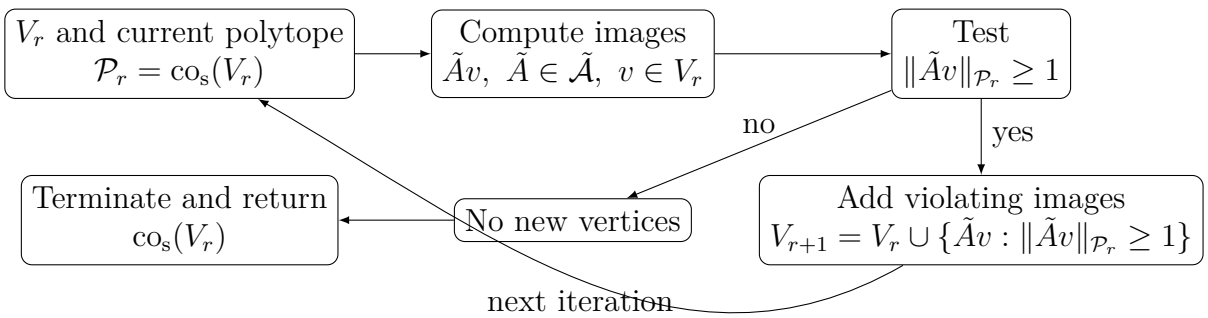

Figure 2.1.: Schematic iteration of the invariant-polytope algorithm. Termination occurs when no new violating images are found.

2.3. Termination Theory

2.3.1. Proof Strategy

The proof path is: (i) introduce invariant and extremal norms, (ii) connect balanced invariant polytopes to extremal norms, and (iii) deduce a valid upper bound for $\mathrm{JSR}(\mathcal{A})$. This provides a finite certificate when the algorithm terminates.

2.3.2. Auxiliary Results

Definition 2.3.1. (Barabanov, 1988) For a family of matrices $\mathcal{A}$ a norm $\|\cdot\|$ is called *invariant* or *Barabanov* if there exists a number $\lambda \geq 0$ such that

$$\max_{A\in\mathcal{A}} \|Ax\| = \lambda\|x\|, \quad \text{for all } x \in \mathbb{R}^d.$$

The scalar in Definition 2.3.1 equals $\mathrm{JSR}(\mathcal{A})$.

Theorem 2.3.2. *(Barabanov, 1988) Every irreducible family $\mathcal{A}$ possesses an invariant norm.*

Example 2.3.3. Consider the family

$$\mathcal{A} = \{A_1, A_2\}, \qquad A_1 = \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix}, \quad A_2 = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}.$$

Both matrices are orthogonal, so the Euclidean norm satisfies

$$\|A_1x\|_2 = \|x\|_2, \qquad \|A_2x\|_2 = \|x\|_2$$

for all $x \in \mathbb{R}^2$. Therefore

$$\max_{A\in\mathcal{A}} \|Ax\|_2 = \|x\|_2$$

for all $x \in \mathbb{R}^2$. Hence $\|\cdot\|_2$ is a Barabanov norm with $\lambda = 1$, and in particular it is extremal for $\mathcal{A}$ with $\mathrm{JSR}(\mathcal{A}) = 1$.

Definition 2.3.4. (Protasov, 1996) A norm $\|\cdot\|$ is called *extremal* for a family of matrices $\mathcal{A}$ if

$$\max_{A\in\mathcal{A}} \|Ax\| \leq \mathrm{JSR}(\mathcal{A})\|x\| \quad \text{for all } x \in \mathbb{R}^d.$$

Theorem 2.3.5. *For every extremal norm $\|\cdot\|$ of a family of matrices $\mathcal{A}$*

$$\max_{P\in\mathcal{A}^k}\|P\| = \mathrm{JSR}(\mathcal{A})^k$$

holds for every $k \in \mathbb{N}$.

Proof. By submultiplicativity and extremality of the norm, we have

$$\begin{aligned}\max_{P\in\mathcal{A}^k}\|P\| &= \max_{d_k,\dots,d_1}\|A_{d_k}\cdots A_{d_1}\| \\ &\le \max_{d_k,\dots,d_1}\|A_{d_k}\|\cdots\|A_{d_1}\| \\ &\le \mathrm{JSR}(\mathcal{A})^k\end{aligned}$$

By the three-member inequality (1.5), we know

$$\begin{aligned}\mathrm{JSR}(\mathcal{A}) &\le \max_{P\in\mathcal{A}^k}\|P\|^{\frac{1}{k}} \\ \mathrm{JSR}(\mathcal{A})^k &\le \max_{P\in\mathcal{A}^k}\|P\|\end{aligned}$$

for every $k \in \mathbb{N}$. □

Remark 2.3.6. In particular, when $k = 1$, we have $\max_{A\in\mathcal{A}}\|A\| = \mathrm{JSR}(\mathcal{A})$. The existence of an extremal norm for $\mathcal{A}$ therefore enables the exact computation of the joint spectral radius.

Definition 2.3.7. A polytope $\mathcal{P}$ is called *balanced* if it spans the whole space and is equal to the symmetrized convex hull of a finite set of vertices.

$$\mathcal{P} = \mathrm{co_s}(V) := \mathrm{co}(V, -V) = \{x \in \mathbb{R}^d : x = \sum_{i=1}^{M}\alpha_i v_i \text{ and } \sum_{i=1}^{M}|\alpha_i| \le 1\}$$

If a polytope satisfies $\mathcal{A}\mathcal{P} \subseteq \mathcal{P}$ for some family $\mathcal{A}$, it is called *invariant under* $\mathcal{A}$.

Definition 2.3.8. (Magaril-Il'yaev and Tikhomirov, 2001) Let $C \subseteq \mathbb{R}^d$ be a nonempty, convex set such that $0 \in \mathrm{int}(C)$. The *Minkowski functional* (or *gauge function*) associated to C is the mapping $p_C : \mathbb{R}^d \to [0, \infty)$ defined by:

$$p_C(x) = \inf\{\lambda > 0 \mid x \in \lambda C\} \quad \text{for all } x \in \mathbb{R}^d.$$

Remark 2.3.9. In general, the Minkowski functional is only a seminorm on $\mathbb{R}^d$. If $\mathcal{P}$ is a balanced polytope, then the associated Minkowski functional defines a norm on $\mathbb{R}^d$. Moreover, every induced matrix norm on $\mathbb{R}^{d\times d}$ is submultiplicative. In both cases, we denote the norm by $\|\cdot\|_{\mathcal{P}}$.

2.3.3. Main Result

Theorem 2.3.10. *(Guglielmi and Protasov, 2013) If a balanced polytope $\mathcal{P}$ is invariant under $\tilde{\mathcal{A}}$ from Section 1.4, the corresponding Minkowski functional is an extremal norm for the family $\mathcal{A}$ and* $\mathrm{JSR}(\mathcal{A}) \leq \hat{\rho}$.

Proof.

$$\begin{aligned} \tilde{\mathcal{A}}\mathcal{P} \subseteq \mathcal{P} &\implies \mathcal{A}\mathcal{P} \subseteq \hat{\rho}\mathcal{P} \\ &\implies \|A\|_{\mathcal{P}} \leq \hat{\rho} \quad \text{for all } A \in \mathcal{A} \\ &\implies \max_{A \in \mathcal{A}} \|A\|_{\mathcal{P}} \leq \hat{\rho}. \end{aligned}$$

□

2.3.4. Termination Conditions

Theorem 2.3.10 gives the central sufficient condition: if the algorithm produces a balanced polytope invariant under $\tilde{\mathcal{A}}$, then the corresponding upper bound is certified.

For finite termination of the implemented vertex-expansion scheme, a necessary algorithmic condition is that only finitely many violating images are generated. If this fails, the vertex set keeps growing and the algorithm does not terminate.

2.4. Implementation Details

2.4.1. Stopping Criterion and Candidate Rejection

The runtime of Algorithm 1 can be high in practice, so it is important to detect incorrect candidates early and control the growth of the vertex set V. A common practical safeguard is to set an upper bound on the number of vertices. For rejecting an incorrect s.m.p. candidate, (Guglielmi and Protasov, 2013) proposed the following method.

In the case where the candidate $\tilde{\Pi}$ has a unique simple leading eigenvalue and there exists only one s.m.p., define v_1^* as the leading eigenvector of $\tilde{\Pi}^*$ (the conjugate operator of $\tilde{\Pi}$), normalized by $(v_1, v_1^*) = 1$, and define $v_i^* := \tilde{A}_{d_i}^* \cdots \tilde{A}_{d_k}^* v_1^*$, which are the leading eigenvectors of the conjugates of the cyclic permutations of $\Pi \in \mathcal{A}^k$.

If $\|Av\|_{\mathrm{cos}(V)} > 1$ and there exists a j such that $|(v_j^*, Av)| > 1$, then the chosen candidate is not an s.m.p. and must be rejected.

2.4.2. Rebalancing and Additional Initial Vertices

Three years after publishing the fundamental algorithm with the stopping criterion, the authors released a new paper on rebalancing multiple s.m.p.s, as well as starting with extra vertices so that the polytope is better conditioned (Guglielmi and Protasov, 2016). In that setting, termination with multiple s.m.p.s can be achieved, but it is substantially more delicate than in the unique-s.m.p. case.

2.5. Remarks

2.5.1. Complex Variants

If the leading eigenvector associated with a candidate product is complex, then the real balanced-polytope norm is no longer sufficient and must be replaced by a *complex-balanced polytope norm*. This changes the convex-hull membership tests in the linear-programming step, while the high-level structure of the algorithm remains the same.

In the present thesis, we restrict attention to the real case. For details on complex variants, we refer to (Guglielmi et al., 2005; Guglielmi and Zennaro, 2007; Guglielmi and Protasov, 2013).

Remark 2.5.1. There are rare cases where the finiteness property does not hold, or where no invariant polytopes exist. In those cases, the invariant-polytope algorithm will not find the value of $\mathrm{JSR}(\mathcal{A})$.

3. Finite-Tree Algorithm

This chapter is a literature-based summary of the finite-tree approach from (Möller and Reif, 2014) and related references. Its role is to provide the theoretical and algorithmic background needed for the hybrid method in Chapter 4.

3.1. Overview

3.1.1. Motivating Ideas

The finite-tree algorithm bounds arbitrary products $P \in \mathcal{A}^k$ by decomposing them into factors that are 1-bounded with respect to a chosen norm.

3.1.2. Problem Formulation

The objective is to construct a finite certificate that implies $\mathrm{JSR}(\mathcal{A}) \leq 1$. Concretely, one seeks a structured subtree such that all relevant leaf products are bounded by 1 and all remaining factors admit polynomial control in their length.

3.1.3. Core Idea and Notation

Definition 3.1.1. A *tree* is a connected, acyclic graph T with a distinguished node t_0 called the *root*. The tree structure induces a natural hierarchy: for any node t_i, there exists a unique path from t_0 to t_i. The tree can then be described in terms of the following concepts:

- **Parent**: For any node $t \in T$, the node preceding t on the path from the root t_0 is called its *parent*.
- **Children**: The nodes adjacent to a given node t that are further from the root t_0 are called its *children*.
- **Leaf**: A node with no children is called a *leaf*.
- **Internal vertex**: A node with at least one child is called an *internal* or *non-leaf* node.

We define $T_{\mathcal{A}}$ as the tree consisting of all products in the set $\bigcup_{k\in\mathbb{N}_0} \mathcal{A}^k$ where every multi-index represents a node. For the empty index ($k = 0$) we set $A_{[]} = I_d$, the identity matrix in $\mathbb{R}^{d\times d}$, and call this node the root.

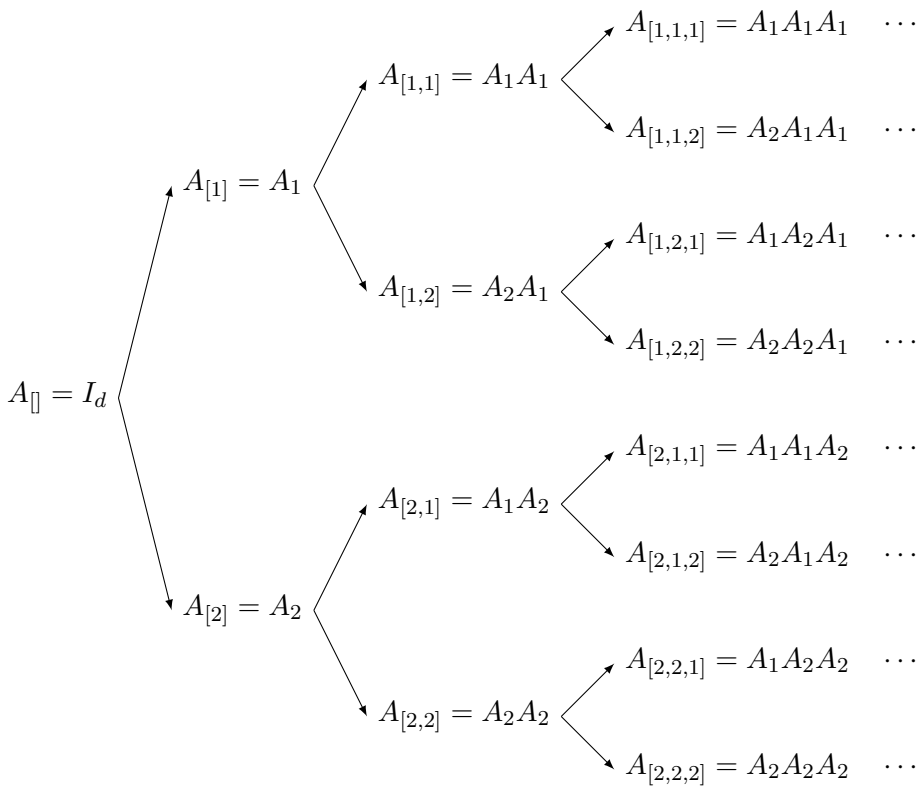

Figure 3.1.: Example of the product tree $T_{\mathcal{A}}$ for $\mathcal{A} = \{A_1, A_2\}$.

In general, a direct search in $T_{\mathcal{A}}$ is not computationally feasible. Following (Möller and Reif, 2014), one introduces generators and set-valued nodes.

We fix a set $\mathcal{G} = \{G_1, \cdots, G_m\}$ with $G_i \in \mathcal{A}^{k_i}$ and $\rho(G_i)^{\frac{1}{k_i}} \leq 1$. All products from this set will be called *generators*. We also define $A_l := G_{-l}$ for $l \in \{-m, \ldots, -1\}$.

Definition 3.1.2. (Möller and Reif, 2014) The set $\mathcal{J}^k := \{-m, \ldots, -1, 1, \ldots, n\}^k$ denotes the collection of all multi-indices of length k with entries in $\{-m, \ldots, -1, 1, \ldots, n\}$. For a multi-index $J = [j_1, \ldots, j_k] \in \mathcal{J}^k$, the corresponding matrix product set is defined by:

$$A_J := \{A_{j_k}^{e_k} \cdots A_{j_1}^{e_1} \mid e_i = 1 \text{ if } j_i > 0 \text{ or } e_i \in \mathbb{N}_0 \text{ otherwise}\}.$$

Definition 3.1.3. Let $J = [j_1, \cdots, j_k]$ be a multi-index and let $j_{k+1} \in \mathcal{J}$. We define the extension

$$[J, j_{k+1}] := [j_1, \cdots, j_k, j_{k+1}].$$

Nested concatenations are always flattened a priori.

Definition 3.1.4. Let $t \in T$ be a node with the multi-index J. A child node $\tilde{t}$ with the multi-index $[J, j_{k+1}]$ is called:

- *positive* if j_{k+1} is positive.
- *generator* if j_{k+1} is negative.

Definition 3.1.5. (Möller and Reif, 2014) A so-called $(\mathcal{A}, \mathcal{G})$-tree T has the following structure:

- The root node contains the identity matrix: $t_0 = \{I_d\}$.
- For every node $t \in T \setminus \{t_0\}$, there is exactly one parent node in T.
- Each node $t \in T$ is either:
  - a leaf,
  - a parent of exactly n positive children,
  - or a parent with at least one generator child.

Definition 3.1.6. (Möller and Reif, 2014) A node in the tree is called *covered* if it is a subset of one of its ancestors in the tree. Otherwise, it is called *uncovered*. The set of uncovered leaves is denoted as

$$\mathcal{L}(T) := \{t \in T : t \text{ is an uncovered leaf}\}$$

and called *leafage* of the tree T.

Definition 3.1.7. (Möller and Reif, 2014) An $(\mathcal{A}, \mathcal{G})$-tree T is called *1-bounded* with respect to a matrix norm $\|\cdot\|$ if

$$\sup_{L \in \mathcal{L}(T)} \|L\| \leq 1.$$

If $\sup_{L \in \mathcal{L}(T)} \|L\| < 1$, then the tree is called *strictly 1-bounded*.

3.2. Algorithmic Formulation

3.2.1. Algorithm Description

The finite-tree algorithm performs a non-exhaustive search in the space of $(\mathcal{A}, \mathcal{G})$-trees. At each active node, it checks whether the corresponding set of products is already 1-bounded. If not, children are generated according to the child-selection rule, and the process is repeated until no active node violates the bound.

3.2.2. Pseudocode

Algorithm 2 Finite-tree algorithm

Input: $\mathcal{A} = \{A_1, \ldots, A_n\}$ and $\mathcal{G} = \{G_1, \cdots, G_m\}$
Output: If the algorithm terminates, the tree condition holds
$Q \leftarrow \{[1], \cdots, [n]\}$
while $Q \neq \emptyset$ **do**
  $Q_{\text{new}} \leftarrow \emptyset$
  for $J \in Q$ **do**
    if $\sup\{\|P\| : P \in A_J\} > 1$ **then**
      $\mathcal{C} \leftarrow \textsc{selectChildren}(J, \mathcal{A}, \mathcal{G})$
      $Q_{\text{new}} \overset{\cup}{\leftarrow} \{[J, j] : j \in \mathcal{C}\}$
  $Q \leftarrow Q_{\text{new}}$

3.2.3. Visualization

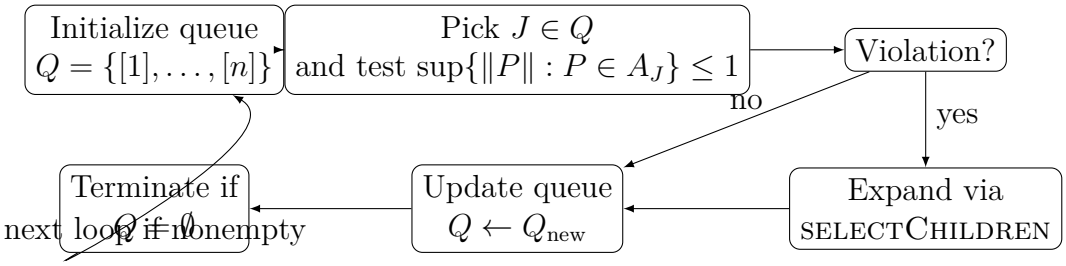

Figure 3.2.: Workflow of the finite-tree algorithm. Termination occurs when no violating node remains in the active queue.

3.3. Termination Theory

3.3.1. Proof Strategy

The proof strategy has two steps. First, one proves a polynomial bound for products encoded by nodes of a fixed $(\mathcal{A}, \mathcal{G})$-tree. Second, one combines this with prefix-leaf decomposition to bound arbitrary products and pass to the k-th root limit in the definition of JSR.

3.3.2. Auxiliary Results

Lemma 3.3.1. *For every $(\mathcal{A}, \mathcal{G})$-tree T there exists a strictly positive, monotone polynomial p such that for every node $t \in T$ and every product $P \in \mathcal{A}^k$ encoded by t, $\|P\| \leq p(k)$ holds.*

Proof. Take an arbitrary node $t \in T$ with multi-index $J = [j_1, \cdots, j_l]$. Now every j_i encodes either a matrix from $\mathcal{A}$ or arbitrary powers of a generator from $\mathcal{G}$. Since all those matrices have spectral radius less than or equal to 1, their Jordan normal forms grow at most polynomially and thus, by norm equivalence in finite-dimensional vector spaces, there exists a strictly positive, monotone polynomial p_i with

$$\|A_{j_i}^e\| \leq p_i(e) \quad \text{for all } e \in \mathbb{N}_0.$$

We define:

$$\begin{aligned} p_t &:= p_l \cdots p_1 \\ p &:= \sum_{t \in T} p_t \end{aligned}$$

The sum over all of T is finite since T has only finitely many nodes. Take now an arbitrary product $P = A_{j_l}^{e_l} \cdots A_{j_1}^{e_1} \in \mathcal{A}^k$ that is encoded by t.

Now we have:

$$\begin{aligned} \|P\| &= \|A_{j_l}^{e_l} \cdots A_{j_1}^{e_1}\| \\ &\leq \|A_{j_l}^{e_l}\| \cdots \|A_{j_1}^{e_1}\| \\ &\leq p_l(e_l) \cdots p_1(e_1) \\ &\leq p_l(k) \cdots p_1(k) \\ &= p_t(k) \leq p(k). \end{aligned}$$

Since the node t and the product P were chosen arbitrary, this concludes the proof. □

3.3.3. Main Result

Theorem 3.3.2. *(Möller and Reif, 2014) If a 1-bounded $(\mathcal{A}, \mathcal{G})$-tree is found for a given matrix set $\mathcal{A}$ and generator set $\mathcal{G}$, then* $\mathrm{JSR}(\mathcal{A}) = 1$.

Proof. We take an arbitrary positive multi-index $I = [i_1, \cdots, i_k]$. Our found $(\mathcal{A}, \mathcal{G})$-tree T allows for a unique decomposition of this multi-index into $I_k, \cdots, I_1$ where

$$A_I = A_{I_k} A_{I_{k-1}} \cdots A_{I_1}.$$

Each sub-multi-index $I_j \quad j = 1, \cdots, k-1$ encodes either a leaf of T or is the empty set and I_k encodes any node of T. Now due to Lemma 3.3.1 there exists a polynomial p

such that

$$\|A_I\| = \|A_{I_k} A_{I_{k-1}} \cdots A_{I_1}\| \leq \|A_{I_k}\| \cdot \|A_{I_{k-1}}\| \cdots \|A_{I_1}\| \leq \|A_{I_k}\| \leq p(k).$$

The last inequality holds since all $I_j \quad j = 1, \cdots, k-1$ are either the empty multi-index or leaves such that their products A_{I_j} are bounded by 1. Furthermore A_{I_k} can be bounded by the polynomial p and the length of the product, which is less than k. Now we have

$$\mathrm{JSR}(\mathcal{A}) = \lim_{k\to\infty} \max_{P\in\mathcal{A}^k} \|P\|^{\frac{1}{k}} \leq \lim_{k\to\infty} p(k)^{\frac{1}{k}} = 1$$

since the chosen multi-index was arbitrary. □

3.3.4. Termination Conditions

Theorem 3.3.2 gives a sufficient condition: finding a 1-bounded $(\mathcal{A}, \mathcal{G})$-tree implies $\mathrm{JSR}(\mathcal{A}) = 1$.

For finite termination of the implemented search procedure, a necessary algorithmic condition is that the child-selection process reaches a finite 1-bounded tree. In practice, this depends on the choice of generators and branching heuristics.

3.4. Implementation Details

3.4.1. Child-Selection Policy

Remark 3.4.1. (Möller and Reif, 2014) The procedure SELECTCHILDREN determines whether all positive children (from 1 to n) or a subset of generator children (from -1 to $-m$) is chosen, according to the theory established there.

3.5. Remarks

Remark 3.5.1. (Möller and Reif, 2014) The decomposition used in the proof of Theorem 3.3.2 exists and is unique due to the structure of an $(\mathcal{A}, \mathcal{G})$-tree.

Remark 3.5.2. (Möller and Reif, 2014) In the cited work, the authors show that the use of generators is essential for ensuring termination in cases where the original invariant-polytope algorithm fails to terminate. Deciding when to use generators instead of all positive children is still part of active research and is handled by heuristics. In particular, termination with multiple s.m.p. candidates is substantially more delicate.

4. Invariant-Tree Algorithm

This chapter contains the thesis-specific contribution: a detailed formulation of the invariant-tree algorithm and a corresponding termination analysis, based on the hybrid framework of (Mejstrik and Reif, 2025). In the recent literature, the same construction is also referred to as the *tree-flavoured-invariant-polytope-algorithm*. Implementation details for finite 1-bounding of generated product sets are taken from the cited work and are not claimed as original contributions here.

4.1. Overview

4.1.1. Combination Strategies

At its core, the invariant-polytope algorithm constructs a polytope norm adapted to the matrix family, while the finite-tree algorithm controls growth through product decompositions. This naturally suggests combining both algorithms.

One option is modular: first run the invariant-polytope algorithm for several iterations, then use the resulting norm inside the finite-tree phase. A second option is concurrent: update norm information and tree-decomposition information in each iteration. Following (Mejstrik and Reif, 2025), we use the concurrent variant.

4.1.2. Problem Formulation

After preprocessing from Section 1.4, we obtain a scaled family $\tilde{\mathcal{A}}$. From the chosen candidate and the three-member inequality (1.5), we already know $\mathrm{JSR}(\tilde{\mathcal{A}}) \geq 1$. Hence proving $\mathrm{JSR}(\tilde{\mathcal{A}}) \leq 1$ is sufficient.

To do this, we show that there exist a polynomial p and a norm $\|\cdot\|$ such that, similarly to Chapter 3, every product can be bounded polynomially in its length:

$$\|P\| \leq p(k) \quad \text{for all } P \in \tilde{\mathcal{A}}^k,\ k \in \mathbb{N}.$$

The polynomial bound is obtained by decomposing each multi-index into finitely many parts arising from 1-v-bounded test trees, plus one residual part controlled by finite-tree arguments.

4.2. Algorithmic Formulation

The hybrid algorithm uses a vertex-wise version of the 1-boundedness concept from Chapter 3.

Definition 4.2.1. Let $v \in \mathbb{R}^d$ and let $\|\cdot\|$ be a vector norm on $\mathbb{R}^d$. An $(\mathcal{A}, \mathcal{G})$-tree T is called 1-*v*-*bounded* with respect to $\|\cdot\|$ if

$$\sup_{L \in \mathcal{L}(T)} \|Lv\| \leq 1.$$

If

$$\sup_{L \in \mathcal{L}(T)} \|Lv\| < 1,$$

then T is called *strictly* 1-*v*-*bounded*.

The algorithm starts, as in the invariant-polytope method, from a set V_0 of starting vertices. Each vertex set V_r defines a vector polytope norm $\|\cdot\|_{V_r}$. At iteration r, for every vertex $v \in V_r$, we build a test $(\mathcal{A}, \mathcal{G})$-tree T_v. If some L in the leafage of T_v satisfies $\|Lv\|_{V_r} \geq 1$, then Lv is added as a new vertex candidate. If all test-tree leafages in one iteration are 1-v-bounded with respect to the current polytope norm, the algorithm terminates.

Algorithm 3 Invariant-tree algorithm

Input: $\mathcal{A}$ (set of matrices), $\mathcal{G}$ (set of generators), V (set of starting vertices)
$V_{\text{new}} \leftarrow V$
while $V_{\text{new}} \neq \emptyset$ **do**
  $V_{\text{rem}} \leftarrow V_{\text{new}}$
  $V_{\text{new}} \leftarrow \emptyset$
  for $v \in V_{\text{rem}}$ **do**
    Construct minimal $(\mathcal{A}, \mathcal{G})$-tree $\mathbf{T}$ as discussed in Section 4.4
    $\mathcal{W} \leftarrow \{Lv : L \in \mathcal{L}(\mathbf{T}) \text{ and } \|Lv\|_{\text{co}_s(V)} \geq 1\}$
    if $\mathcal{W} \neq \emptyset$ **then**
      $V \overset{\cup}{\leftarrow} \mathcal{W}$
      $V_{\text{new}} \overset{\cup}{\leftarrow} \mathcal{W}$

4.3. Termination Theory

Theorem 4.3.1. *If Algorithm 3 terminates, then* $\text{JSR}(\mathcal{A}) \leq 1$.

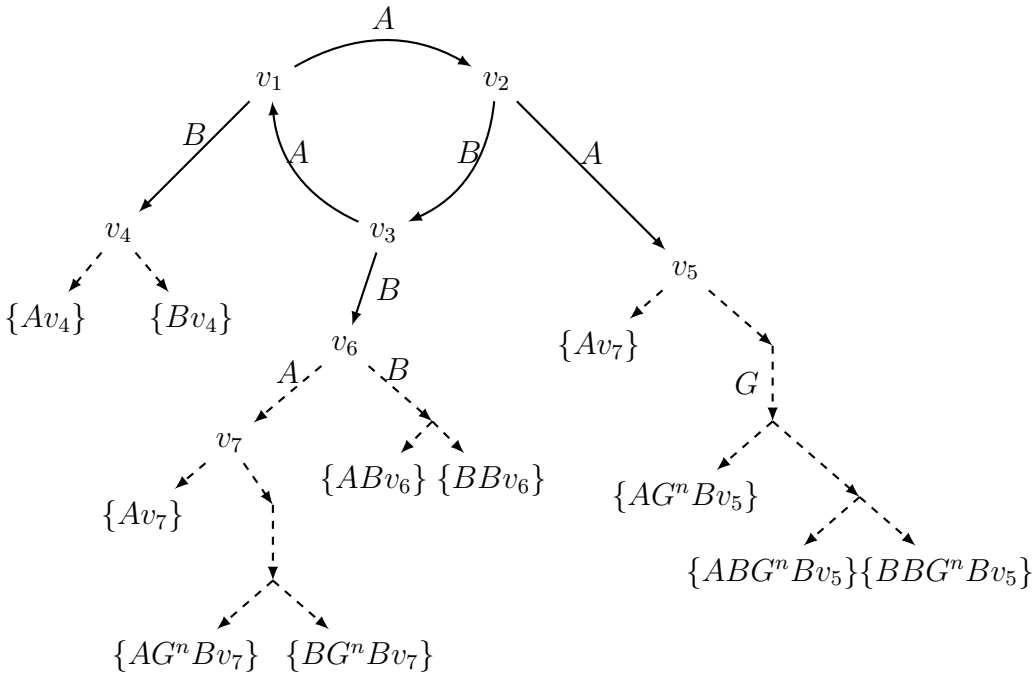

Figure 4.1.: Cyclic tree structure generated by the algorithm. Here $\mathcal{A}$ consists of two matrices A and B and the chosen candidate is $\Pi = ABA$. Dashed arrows indicate $(\mathcal{A}, \mathcal{G})$-trees used for bounds. Solid arrows indicate vertices added to V in one iteration. Leaf labels with powers G^n represent matrix-vector products that remain 1-v-bounded for all $n \in \mathbb{N}_0$.

4.3.1. Proof Strategy

Products are bounded by norms built during intermediate iterations, and submultiplicativity allows us to pass to the final norm. Therefore, we work with $\|\cdot\|_V$ for the final vertex set $V = \cup_{i=1}^{K} V_i$, where K is the final iteration. The proof generalizes the finite-tree argument by repeatedly expanding linear combinations and controlling their coefficients via 1-v-boundedness.

4.3.2. Auxiliary Termination Results

Lemma 4.3.2. *If Algorithm 3 terminates, there exists a monotone polynomial p such that for every product $P \in \mathcal{A}^k$ and $v \in V$, $\|Pv\|_{co_s(V)} \leq p(k)$ holds.*

Proof. If the algorithm terminates, then the set of vertices $V = \{v_1, \cdots, v_m\}$ is finite and for each $v \in V$ there exists an $(\mathcal{A}, \mathcal{G})$-tree T_v that is 1-v-bounded with respect to $\|\cdot\|_{\mathrm{co_s}(V)}$. In particular, for every $L \in \mathcal{L}(T_v)$ one has $\|Lv\|_{\mathrm{co_s}(V)} \leq 1$, which means there exists $\lambda_1, \cdots, \lambda_m$ such that $Lv = \sum \lambda_j v_j$ with $\sum |\lambda_j| \leq 1$. According to Lemma 3.3.1, there

exists a polynomial p_v such that for every product $P \in \mathcal{A}^k$ encoded by T_v, $\|Pv\|_{\mathrm{cos}(V)} \leq p_v(k)$. Take an arbitrary product $P \in \mathcal{A}^k$ and $v \in V$. Define:

$$\begin{aligned}
p &:= \sum_{v \in V} p_v \\
P'L &:= P \text{ with } L \in \mathcal{L}(T_v) \text{ if } P \notin T_v \text{ and } L := I \text{ else} \\
P'_{i_1 \ldots i_r} L_{i_1 \ldots i_r} &:= P'_{i_1 \ldots i_{r-1}} \text{ with } L_{i_1 \ldots i_r} \in \mathcal{L}(T_{v_{i_r}}) \text{ if } P'_{i_1 \ldots i_{r-1}} \notin T_{v_{i_r}} \\
\text{and } L_{i_1 \ldots i_r} &:= I \text{ else.} \\
\text{Where } i_j \in \mathcal{I} &:= \{1, \cdots, m\}
\end{aligned}$$

These decompositions exist and are unique because of the special structure of $(\mathcal{A}, \mathcal{G})$-trees; details can be found in (Möller and Reif, 2014).
Now we have:

$$\begin{aligned}
\|Pv\|_{cos(V)} &= \|P'Lv\|_{cos(V)} \\
&= \|P' \sum_{i_1 \in \mathcal{I}} \lambda_{i_1} v_{i_1}\|_{cos(V)} \\
&= \sum_{i_1 \in \mathcal{I}} |\lambda_{i_1}| \|P' v_{i_1}\|_{cos(V)} \\
&= \sum_{i_1 \in \mathcal{I}} |\lambda_{i_1}| \|P'_{i_1} L_{i_1} v_{i_1}\|_{cos(V)} \\
&= \sum_{i \in \mathcal{I}^2} |\lambda_{i_1}| |\lambda_{i_2}| \|P'_i v_{i_2}\|_{cos(V)} \\
&\ldots \\
&= \sum_{i \in \mathcal{I}^k} |\lambda_{i_1}| \cdot \ldots \cdot |\lambda_{i_k}| \|P'_i v_{i_k}\|_{cos(V)} \\
&\leq \sum_{i \in \mathcal{I}^k} |\lambda_{i_1}| \cdot \ldots \cdot |\lambda_{i_k}| \; p_{v_{i_k}}(k) \\
&\leq p(k) \sum_{i \in \mathcal{I}^k} |\lambda_{i_1}| \cdot \ldots \cdot |\lambda_{i_k}| \\
&= p(k) \sum_{i_1 \in \mathcal{I}} |\lambda_{i_1}| \sum_{i_2 \in \mathcal{I}} |\lambda_{i_2}| \ldots \sum_{i_k \in \mathcal{I}} |\lambda_{i_k}| \\
&\leq p(k) \cdot 1 \cdot \ldots \cdot 1 \\
&\leq p(k)
\end{aligned}$$

Here the first inequality works since after k steps the prefix P'_i is guaranteed to be in the tree $T_{v_{i_k}}$, and thus the norm is bounded by $p_{v_{i_k}}(k)$, which is less than $p(k)$ by definition. The second-to-last inequality follows because the sums of absolute coefficients in the corresponding linear combinations are bounded by 1 due to 1-v-boundedness of the $(\mathcal{A}, \mathcal{G})$-trees. □

4.3.3. Main Termination Result

Proof. Let $I \in \mathcal{I}^k$ be a multi-index of length k. We want to find a polynomial p such that $\|A_I\| \leq p(k)$. We make the assumption all products $P' \in \mathcal{A}^{k'}$ with length k' less than k are already bounded by p i.e.

$$\|P'\| \leq p(k')$$

in their lengths k'. Now by Dicomp-Lemma we get:

$$\begin{aligned}\|A_I v\| &= \|A_{P_{I,v}} A_{L_{I,v}} v\| = \|A_{P_{I,v}}(\lambda_1 v_1 + ... + \lambda_l v_l)\| \\ &= \lambda_1 \|A_{P_{I,v}} v_1\| + ... + \lambda_1 \|A_{P_{I,v}} v_1\| \\ &\leq p(k)\end{aligned}$$

□

Corollary 4.3.3. *If the chosen trees are always $T_v = \mathcal{A}$, then Algorithm 3 reduces to the invariant-polytope algorithm 1. Hence, the invariant-tree algorithm generalizes the invariant-polytope algorithm.*

Corollary 4.3.4. *Assume Algorithm 3 terminates and all test trees T_v are chosen without generator children. Then the polynomial bound in Lemma 4.3.2 can be chosen constant, i.e. there exists $C > 0$ such that*

$$\|P\|_{\mathrm{co_s}(V)} \leq C \quad \textit{for all } P \in \mathcal{A}^k,\ k \in \mathbb{N}.$$

Consequently, for the rescaled family $\widehat{\mathcal{A}} := C^{-1}\mathcal{A}$, one has

$$\widehat{\mathcal{A}}\,\mathrm{co_s}(V) \subseteq \mathrm{co_s}(V),$$

so polytope invariance is recovered after scaling.

Corollary 4.3.5. *Assume Algorithm 3 terminates and all test trees coincide, i.e. $T_v = T$ for all $v \in V$. Then the product bound can be derived via a single prefix-leaf decomposition (PL decomposition) in matrix norm, as in Theorem 3.3.2, without vertex-wise recursive decompositions.*

More precisely, if

$$A_I = A_{I_r} \cdots A_{I_1}$$

is the PL decomposition induced by T, then

$$\|A_I\|_{\mathrm{co_s}(V)} \leq \|A_{I_r}\|_{\mathrm{co_s}(V)} \prod_{j=1}^{r-1} \|A_{I_j}\|_{\mathrm{co_s}(V)} \leq \|A_{I_r}\|_{\mathrm{co_s}(V)} \leq p(k),$$

where p is the polynomial from Lemma 3.3.1.

4.4. Implementation Considerations

Definition 4.4.1. An $(\mathcal{A}, \mathcal{G})$-tree with one generator branch at the root, one path from the root to the covered node, and all branches required to satisfy Definition 3.1.5, is called a *minimal tree*. This tree depends only on $\mathcal{A}$ and the chosen generator G. It is denoted by $T_{(\mathcal{A},\mathcal{G})}$ and, when $\mathcal{A}$ is clear from context, shortened to $T_{\mathcal{G}}$.

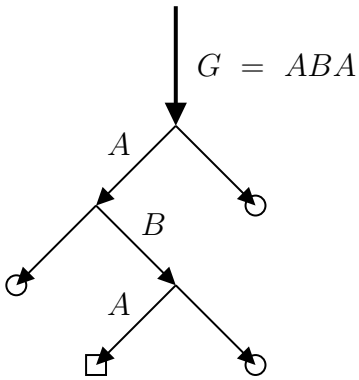

Figure 4.2.: Minimal tree structure $T_{(G,\mathcal{A})}$ for the candidate $G = ABA$ and $\mathcal{A} = \{A, B\}$. The arrows marked with circles are uncovered-leafs, the arrow marked with a square is a covered-leaf

In principle, the invariant-tree algorithm allows a different test tree for each vertex $v \in V$. More precisely, both the test tree and the corresponding bounding tree can be chosen independently for each vertex, as long as they satisfy the $(\mathcal{A}, \mathcal{G})$-tree conditions from Definition 3.1.5.

In our implementation, we use minimal trees only. This means that for each vertex we choose a tree that keeps the generator branch and the first covered-node effect, while minimizing the size of leafage to be tested. This reduces the number of norm checks $\|Lv\|_{\mathrm{cos}(V)}$ in each iteration and is the default tree-selection policy.

Because generator branches encode families such as $\{G^r : r \in \mathbb{N}\}$, a naive implementation would require infinitely many tests. Following Section 3.3 of (Mejstrik and Reif, 2025), we split the generator dynamics into a leading-eigenspace part and a stable remainder, estimate the remainder by pseudo-spectral-radius bounds, and then reduce each infinite branch check to finitely many verifiable conditions for the current vertex.

The practical implementation also uses a maximal iteration counter $r_{\max}$. If no certificate is found before reaching $r_{\max}$, the algorithm stops and returns an approximation interval for $\mathrm{JSR}(\mathcal{A})$ via the three-member inequality:

$$\underline{\rho} \;\leq\; \mathrm{JSR}(\mathcal{A}) \;\leq\; \overline{\rho}.$$

In practice, these bounds can already be accurate even when full termination is not achieved.

From a computational perspective, runtime is dominated by norm/membership checks over leafage elements and by tree construction. Memory is dominated by the current vertex set V and temporary tree data. Restricting to minimal trees keeps leafage small, so both runtime and memory growth are significantly lower than with larger test trees.

Finally, candidate handling from the pure invariant-polytope algorithm is still used in the invariant-tree setting. In particular, candidate rejection and candidate updates can be applied before or during invariant-tree iterations to improve robustness.

4.5. Remarks

4.5.1. Polytope Invariance

After Algorithm 3 terminates, the polytope $\mathrm{co_s}(V)$ is, in general, not invariant under the full family $\mathcal{A}$. This is expected, since the invariant-tree algorithm enforces polynomial product bounds rather than one-step invariance under each matrix in $\mathcal{A}$.

The termination certificate is first vertex-wise. For every $v \in V$, the algorithm provides a test tree T_v such that each leaf matrix in $\mathcal{L}(T_v)$ maps v back into the polytope. Equivalently, with

$$\Gamma_V := \{(v, L) : v \in V,\ L \in \mathcal{L}(T_v)\},$$

one has

$$(v, L) \in \Gamma_V \quad \Longrightarrow \quad Lv \in \mathrm{co_s}(V).$$

A stronger, uniform statement is obtained if one can construct a single $(\mathcal{A}, \mathcal{G})$-tree $T_\star$ that bounds all vertices simultaneously, i.e.

$$\|Lv\|_{\mathrm{co_s}(V)} \leq 1 \quad \text{for all } v \in V \text{ and } L \in \mathcal{L}(T_\star).$$

Then $T_\star$ induces leafage invariance on the whole polytope:

$$x \in \mathrm{co_s}(V),\ L \in \mathcal{L}(T_\star) \quad \Longrightarrow \quad Lx \in \mathrm{co_s}(V).$$

Indeed, writing $x = \sum_{i=1}^{M} \alpha_i v_i$ with $\sum_{i=1}^{M} |\alpha_i| \leq 1$ and $v_i \in V$, one gets

$$Lx = \sum_{i=1}^{M} \alpha_i L v_i,$$

and each term Lv_i belongs to $\mathrm{co_s}(V)$ by assumption. By convexity and symmetry of $\mathrm{co_s}(V)$, this implies $Lx \in \mathrm{co_s}(V)$.

The existence and efficient construction of such a global tree $T_\star$ from the family $(T_v)_{v \in V}$ is a separate algorithmic problem and remains part of future work.

This invariance interpretation, together with the intended play on words relative to the invariant-polytope algorithm, motivates the name *invariant-tree algorithm.*

4.5.2. Memory and Time Complexity

We summarize the computational complexity in terms of the final vertex set

$$V = \{v_1, \dots, v_M\}, \qquad M := |V|.$$

By Algorithm 3, each vertex is processed once after it is added to V. For a processed vertex v, write

$$\ell_v := |\mathcal{L}(T_v)|$$

for the size of the tested leafage, and denote by $C_{\text{tree}}(v)$ the cost of constructing T_v.

For each leaf matrix $L \in \mathcal{L}(T_v)$, one needs to evaluate $\|Lv\|_{\text{co}_s(V)}$, which consists of

$$C_{\text{mv}}(d) \text{ (matrix-vector product)} \quad + \quad C_{\text{norm}}(d, |V|) \text{ (polytope-norm / membership test).}$$

Hence the total runtime is bounded by

$$T_{\text{total}} \lesssim \sum_{v \in V} \Big(C_{\text{tree}}(v) + \ell_v \Big(C_{\text{mv}}(d) + C_{\text{norm}}(d, |V|) \Big) \Big).$$

With uniform bounds $C_{\text{tree}}(v) \le C_{\text{tree}}^{\max}$ and $\ell_v \le \ell_{\max}$, this yields

$$T_{\text{total}} = \mathcal{O}\Big(M \Big(C_{\text{tree}}^{\max} + \ell_{\max} \Big(C_{\text{mv}}(d) + C_{\text{norm}}(d, M) \Big) \Big) \Big).$$

Generator branches introduce an additional preprocessing/verification cost (Section 3.3 in (Mejstrik and Reif, 2025)), where infinite families are reduced to finitely many checks via spectral estimates and pseudo-spectral-radius bounds. This overhead is typically amortized across vertices when the same generators are reused.

The peak memory footprint has the structure

$$S_{\text{total}} \lesssim \underbrace{\mathcal{O}(nd^2)}_{\text{input matrices}} + \underbrace{\mathcal{O}(Md)}_{\text{vertex set}} + \underbrace{\mathcal{O}(S_{\text{tree}}^{\max})}_{\text{active tree data}} + \underbrace{\mathcal{O}(S_{\text{norm}}(d, M))}_{\text{norm/membership solver state}},$$

where $S_{\text{tree}}^{\max}$ denotes the maximal storage needed for one active test tree.

Practically, minimal trees reduce ℓ_v and $S_{\text{tree}}^{\max}$ substantially; if all test trees coincide, tree construction and PL-decomposition data can be reused globally (Corollary above), reducing both runtime and memory constants. Nevertheless, no global polynomial-time bound in (n, d) is expected in full generality, consistent with known hardness results for JSR computation (Tsitsiklis and Blondel, 1997; Blondel and Tsitsiklis, 2000).

5. Numerics

This chapter compares the invariant-polytope algorithm (Algorithm 1), the finite-tree algorithm (Algorithm 2), and the invariant-tree algorithm (Algorithm 3) with respect to iterations, runtime, and solution quality.

5.1. Solves on generated Matrix Sets

In this section we compare all three algorithms on generated matrix families across increasing dimensions. The focus is quantitative scalability with respect to dimension and number of tested instances.

5.2. Solved Problems with One s.m.p.

This section reports comparisons on benchmark families with a unique s.m.p., with iteration count (steps to termination) as the primary speed criterion.

5.3. Solved Problems with Multiple s.m.p.s

This section compares the effective solve space for multiple-s.m.p. instances, i.e., which classes of problems can be solved reliably and with what certificate quality by the new algorithm.

6. Conclusions and Outlook

6.1. Summary of Invariant-Tree numerics

6.2. Resolved Problem Classes

6.3. Remaining Open Problem Classes

6.4. Directions for Future Research

6.4.1. Optimization of Tree Choice and Vertex Selection

6.4.2. Extension to the Complex Setting

6.4.3. Optimization for Nonnegative Matrix-Sets

A. Notation Glossary

A.1. Constants and Globally Defined Variables

A_i Matrix in the family $\mathcal{A}$.

d Dimension of each matrix $A_i \in \mathbb{R}^{d \times d}$.

n Number of matrices in $\mathcal{A}$.

$\mathcal{A} = \{A_1, \ldots, A_n\}$ Base set of matrices.

$\tilde{\mathcal{A}} = \{\tilde{A}_1, \ldots, \tilde{A}_n\}$ Scaled matrix set after preprocessing.

m Number of generators.

$\mathcal{G} = \{G_1, \ldots, G_m\}$ Set of generators.

k Length of a product.

α Scalar.

λ Eigenvalue.

x Vector.

ρ Spectral radius.

σ Singular value.

JSR Joint spectral radius operator.

A.2. Products and Multi-Indices

$\mathcal{I} = \{1, \ldots, n\}$ Positive matrix-index set.

$\mathcal{I}^k$ Set of positive multi-indices of length k.

i Positive index.

I Positive multi-index of length k, i.e. $I = [i_1, \ldots, i_k] \in \mathcal{I}^k$.

$A_I = A_{[i_1,\ldots,i_k]} = A_{i_k} \cdots A_{i_1}$ Product encoded by I.

$\mathcal{J} = \{-m, \dots, -1, 1, \dots, n\}$ Extended index set including generator indices.

$\mathcal{J}^k$ Set of extended multi-indices of length k.

j Extended index.

$\mathcal{A}_\mathcal{G}$ Family of primitive product sets:

$$\mathcal{A}_\mathcal{G} := \{\{G_m^r : r \in \mathbb{N}\}, \dots, \{G_1^r : r \in \mathbb{N}\}, \{A_1\}, \dots, \{A_n\}\}\,.$$

$A_j \in \mathcal{A}_\mathcal{G}$ Primitive set encoded by index j.

J Extended multi-index of length k, i.e. $J = [j_1, \dots, j_k] \in \mathcal{J}^k$.

$A_J = A_{[j_1,\dots,j_k]}$ Product set encoded by J.

$[J_1, j_3, J_4, [j_1, J_2]]$ Expressions are interpreted in flattened form.

A.3. Norms and JSR

$\|\cdot\|_{\mathrm{co_s}(V)} = \|\cdot\|_V$ Polytope norm notation used for matrices and vectors.

$\mathrm{JSR}(\mathcal{A}) := \lim\limits_{k\to\infty} \max\limits_{P\in\mathcal{A}^k} \|P\|^{1/k}$ Norm-based definition of JSR.

$\mathrm{JSR}(\mathcal{A}) = \limsup\limits_{k\to\infty} \max\limits_{P\in\mathcal{A}^k} \rho(P)^{1/k}$ Berger–Wang characterization.

A.4. Algorithmic Notation

Invariant-polytope algorithm Baseline algorithm from Chapter 2.

Finite-tree algorithm Tree-based algorithm from Chapter 3.

Invariant-tree algorithm Combined algorithm from Chapter 4.

M Number of vertices in a finite vertex set representation $V = \{v_1, \dots, v_M\}$.

V_k Set of vertices at step k.

R_k Set of vertices remaining to test at step k.

V Union of vertex sets over all steps.

$\overset{\cup}{\leftarrow}$ Appending by union in pseudocode.

A.5. Tree Terminology

T a tree is a directed graph.

Root Node without a predecessor.

Child and parent If a node $t \in T$ has predecessor $p \in T$, then t is a child of p and p is the parent of t.

Leaf Node without children.

Internal node Node with at least one child.

$T_{\mathcal{A}}$ Infinite product tree generated by $\mathcal{A}$.

$(\mathcal{A}, \mathcal{G})$**-tree** Set-valued tree with root $\{I_d\}$ and internal nodes having either n positive children or at least one generator child.

Covered or uncovered node A node is covered if its set is a subset of an ancestor node; otherwise it is uncovered.

$\mathcal{L}(T)$ Leafage of T, i.e. all uncovered leaves.

1**-bounded** All matrix sets encoded by $\mathcal{L}(T)$ are bounded by 1 in a chosen matrix norm.

1-v**-bounded** All images of matrices encoded by $\mathcal{L}(T)$ applied to a fixed vector v are bounded by 1 in a chosen vector norm.

PL decomposition Prefix-leaf decomposition.

$T_{\min(\mathcal{A},\mathcal{G})}$ Minimal $(\mathcal{A}, \mathcal{G})$-tree, i.e. the smallest such tree containing a covered node.

Bibliography

Alizadeh, F. and Goldfarb, D. (2003). Second-order cone programming. *Mathematical Programming*, 95(1):3–51.

Andersen, E. D., Roos, C., and Terlaky, T. (2003). On implementing a primal-dual interior-point method for conic quadratic optimization. *Mathematical Programming*, 95(2):249–277.

Ando, T. and Shih, M.-H. (1998). Simultaneous contractibility. *SIAM Journal on Matrix Analysis and Applications*, 19(2):487–498.

Barabanov, N. E. (1988). Lyapunov indicator for discrete inclusions, i–iii. *Automation and Remote Control*, 49(2):152–157.

Berger, M. A. and Wang, Y. (1992). Bounded semigroups of matrices. *Linear Algebra and its Applications*, 166:21–27.

Berstel, J. (2005). Growth of repetition-free words—a review. *Theoretical Computer Science*, 340(2):280–290.

Blondel, V. D., Jungers, R. M., and Protasov, V. Y. (2006). On the complexity of computing the capacity of codes that avoid forbidden difference patterns. *IEEE Transactions on Information Theory*, 52(11):5122–5127.

Blondel, V. D. and Nesterov, Y. (2005). Computationally efficient approximations of the joint spectral radius. *SIAM Journal on Matrix Analysis and Applications*, 27(1):256–272.

Blondel, V. D., Nesterov, Y., and Theys, J. (2005). On the accuracy of the ellipsoid norm approximation of the joint spectral radius. *Linear Algebra and its Applications*, 394:91–107.

Blondel, V. D., Theys, J., and Vladimirov, A. A. (2003). An elementary counterexample to the finiteness conjecture. *SIAM Journal on Matrix Analysis and Applications*, 24(4):963–970.

Blondel, V. D. and Tsitsiklis, J. N. (2000). Approximating the spectral radius of sets of matrices in the max-algebra is NP-hard. *IEEE Transactions on Automatic Control*, 45(9):1762–1765.

Cassaigne, J. (1993). Counting overlap-free binary words. In *STACS 93*, volume 665 of *Lecture Notes in Computer Science*, pages 216–225, Berlin. Springer.

Cavaretta, A. S., Dahmen, W., and Micchelli, C. A. (1991). *Stationary Subdivision*. Number 453 in Memoirs of the American Mathematical Society. American Mathematical Society.

Cicone, A., Guglielmi, N., Serra-Capizzano, S., and Zennaro, M. (2010). Finiteness property of pairs of 2
*times*2 sign-matrices via real extremal polytope norms. *Linear Algebra and its Applications*, 432(2–3):796–816.

Daubechies, I. and Lagarias, J. C. (1992). Two-scale difference equations. II. local regularity, infinite products of matrices and fractals. *SIAM Journal on Mathematical Analysis*, 23(4):1031–1079.

Elsner, L. (1995). The generalized spectral-radius theorem: an analytic-geometric proof. *Linear Algebra and its Applications*, 220:151–159.

Finch, S. (2008). Private communication. Private communication.

Finch, S., Sebah, P., and Bai, Z.-Q. (2008). Odd entries in pascal's trinomial triangle.

Fornasini, E. and Valcher, M. E. (2010). Stabilizability of discrete-time positive switched systems. In *Proceedings of the 49th IEEE Conference on Decision and Control (CDC 2010)*, pages 432–437.

Goldwasser, J., Klostermeyer, W., Mays, M., and Trapp, G. (1999). The density of ones in pascal's rhombus. *Discrete Mathematics*, 204:231–236.

Golub, G. H. and Van Loan, C. F. (2013). *Matrix Computations*. Johns Hopkins University Press, Baltimore, MD.

Gripenberg, G. (1996). Computing the joint spectral radius. *Linear Algebra and its Applications*, 234:43–60.

Guglielmi, N., Manni, C., and Vitale, D. (2011). Convergence analysis of c^2 hermite interpolatory subdivision schemes by explicit joint spectral radius formulas. *Linear Algebra and its Applications*, 434:784–902.

Guglielmi, N. and Protasov, V. Y. (2013). Exact computation of joint spectral characteristics of linear operators. *Foundations of Computational Mathematics*, 13(1):37–97.

Guglielmi, N. and Protasov, V. Y. (2016). Invariant polytopes of linear operators with applications to regularity of wavelets and of subdivisions. *SIAM Journal on Matrix Analysis and Applications*, 37(1):18–52.

Guglielmi, N. and Protasov, V. Y. (2024). Computing the spectral gap of a family of matrices. *Mathematics of Computation*, 93(345):259–291.

Guglielmi, N., Wirth, F., and Zennaro, M. (2005). Complex polytope extremality results for families of matrices. *SIAM Journal on Matrix Analysis and Applications*, 27(3):721–743.

Guglielmi, N. and Zennaro, M. (2007). Balanced complex polytopes and related vector and matrix norms. *Journal of Convex Analysis*, 14:729–766.

Guglielmi, N. and Zennaro, M. (2008). An algorithm for finding extremal polytope norms of matrix families. *Linear Algebra and its Applications*, 428(10):2265–2282.

Guglielmi, N. and Zennaro, M. (2009). Finding extremal complex polytope norms for families of real matrices. *SIAM Journal on Matrix Analysis and Applications*, 31(2):602–620.

Gurvits, L. (1995). Stability of discrete linear inclusions. *Linear Algebra and its Applications*, 231:47–85.

Hare, K. G., Morris, I. D., Sidorov, N., and Theys, J. (2011). An explicit counterexample to the lagarias–wang finiteness conjecture. *Advances in Mathematics*, 226(6):4667–4701.

Hassan, M. F., Ivrissimitzis, I. P., Dodgson, N. A., and Sabin, M. A. (2002). An interpolating 4-point c^2 ternary stationary subdivision scheme. *Computer Aided Geometric Design*, 19(1):1–18.

Hechler, J., Mößner, B., and Reif, U. (2009). c^1-continuity of the generalized four-point scheme. *Linear Algebra and its Applications*, 430(11–12):3019–3029.

Horn, R. A. and Johnson, C. R. (1985). *Matrix Analysis*. Cambridge University Press, Cambridge.

Jungers, R. M. (2009). *The Joint Spectral Radius: Theory and Applications*, volume 385 of *Lecture Notes in Control and Information Sciences*. Springer, Berlin.

Jungers, R. M. and Protasov, V. Y. (2009). Counterexamples to the complex polytope extremality conjecture. *SIAM Journal on Matrix Analysis and Applications*, 31(2):404–409. Preprint: arXiv:0802.2654.

Jungers, R. M., Protasov, V. Y., and Blondel, V. D. (2008). Efficient algorithms for deciding the type of growth of products of integer matrices. *Linear Algebra and its Applications*, 428(10):2296–2312.

Jungers, R. M., Protasov, V. Y., and Blondel, V. D. (2009). Overlap-free words and spectra of matrices. *Theoretical Computer Science*, 410(38–40):3670–3684.

Kozyakin, V. (2005). A dynamical systems construction of a counterexample to the finiteness conjecture. In *Proceedings of the 44th IEEE Conference on Decision and Control and ECC 2005*, pages 2338–2343.

Kozyakin, V. S. (1990). Algebraic unsolvability of problem of absolute stability of desynchronized systems. *Automation and Remote Control*, 51(6):754–759.

Kozyakin, V. S. (2009). On the computational aspects of the theory of joint spectral radius. *Doklady Mathematics*, 80(1):487–491.

Lagarias, J. C. and Wang, Y. (1995). The finiteness conjecture for the generalized spectral radius of a set of matrices. *Linear Algebra and its Applications*, 214.

Maesumi, M. (1996). An efficient lower bound for the generalized spectral radius. *Linear Algebra and its Applications*, 240:1–7.

Maesumi, M. (2000). Joint spectral radius and hölder regularity of wavelets. *Computers & Mathematics with Applications*, 40(1):145–155.

Maesumi, M. (2008). Optimal norms and the computation of joint spectral radius of matrices. *Linear Algebra and its Applications*, 428(10):2324–2338.

Magaril-Il'yaev, G. G. and Tikhomirov, V. M. (2001). *Convex Analysis: Theory and Applications*, volume 222 of *Translations of Mathematical Monographs*. American Mathematical Society. Translated from Russian by Dmitry Chibisov.

Mason, O. and Shorten, R. N. (2007). Quadratic and copositive lyapunov functions and the stability of positive switched linear systems. In *Proceedings of the 2007 American Control Conference (ACC 2007)*, pages 657–662, New York City, NY, USA. IEEE. Marriott Marquis Hotel at Times Square, July 11–13, 2007.

Mejstrik, T. (2020). Improved invariant polytope algorithm and applications. *ACM Transactions on Mathematical Software*, 46(3):29:1–29:26.

Mejstrik, T. and Reif, U. (2025). A hybrid approach to joint spectral radius computation. *Linear Algebra and its Applications*.

Mengi, E. (2006). *Measures for Robust Stability and Controllability*. PhD thesis, New York University.

Moision, B. E., Orlitsky, A., and Siegel, P. H. (2001). On codes that avoid specified differences. *IEEE Transactions on Information Theory*, 47(1):433–442.

Möller, C. (2015). *A New Strategy for Exact Determination of the Joint Spectral Radius*. PhD thesis, TU Darmstadt.

Möller, C. and Reif, U. (2014). A tree-based approach to joint spectral radius determination. *Linear Algebra and its Applications*, 463:154–170.

Parrilo, P. A. and Jadbabaie, A. (2008). Approximation of the joint spectral radius using sum of squares. In *Linear Algebra and its Applications*, volume 428, pages 2385–2402.

Protasov, V. Y. (1996). The joint spectral radius and invariant sets of linear operators. *Fundamentalnaya i Prikladnaya Matematika*, 2(1):205–231.

Protasov, V. Y. (1997). The generalized spectral radius. a geometric approach. *Izvestiya: Mathematics*, 61(5):995–1030.

Protasov, V. Y. (2000). Asymptotic behaviour of the partition function. *Sbornik: Mathematics*, 191(3–4):381–414.

Protasov, V. Y. (2004). On the regularity of de rham curves. *Izvestiya: Mathematics*, 68(3):27–68.

Protasov, V. Y. (2010). Invariant functionals for random matrices. *Functional Analysis and Its Applications*, 44(3):230–233.

Protasov, V. Y., Jungers, R. M., and Blondel, V. D. (2010). Joint spectral characteristics of matrices: a conic programming approach. *SIAM Journal on Matrix Analysis and Applications*, 31(4):2146–2162.

Reznick, B. (1990). Some binary partition functions. In *Analytic Number Theory: Proceedings of a Conference in Honor of P.T. Bateman*, pages 451–477. Birkhäuser, Boston.

Rota, G.-C. and Strang, G. (1960). A note on the joint spectral radius. *Indagationes Mathematicae*, 22(4):379–381.

Rudin, W. (1976). *Principles of Mathematical Analysis*. McGraw-Hill, New York, 3 edition.

Stewart, G. W. and Sun, J. G. (1990). *Matrix Perturbation Theory*. Academic Press, New York.

Strang, G. (2001). The joint spectral radius (commentary). In *Collected Works of Gian-Carlo Rota*. Birkhäuser, Basel. Commentary on Rota–Strang (1960).

Tsitsiklis, J. N. and Blondel, V. D. (1997). The lyapunov exponent and joint spectral radius of pairs of matrices are hard—when not impossible—to compute and to approximate. *Mathematics of Control, Signals and Systems*, 10(1):31–40.

Vagnoni, C. and Zennaro, M. (2009). The analysis and the representation of balanced complex polytopes in 2D. *Foundations of Computational Mathematics*, 9(3):259–294.

Vandergraft, J. S. (1968). Spectral properties of matrices which have invariant cones. *SIAM Journal on Applied Mathematics*, 16:1208–1222.

Villemoes, L. (1994). Wavelet analysis of refinement equations. *SIAM Journal on Mathematical Analysis*, 25(5):1433–1460.

Zheng, H., Zhao, H., Ye, Z., and Zhou, M. (2007). Differentiability of a 4-point ternary subdivision scheme and its application. *IAENG International Journal of Applied Mathematics*, 36(1):19–24.